

CancerDiagnosis

August 4, 2019

Personalized cancer diagnosis

1. Business Problem

1.1. Description

Source: <https://www.kaggle.com/c/msk-redefining-cancer-treatment/>

Data: Memorial Sloan Kettering Cancer Center (MSKCC)

Download training_variants.zip and training_text.zip from Kaggle.

Context:

Source: <https://www.kaggle.com/c/msk-redefining-cancer-treatment/discussion/35336#198462>

Problem statement :

Classify the given genetic variations/mutations based on evidence from text-based clinical literature.

1.2. Source/Useful Links

Some articles and reference blogs about the problem statement

1. <https://www.forbes.com/sites/matthewherper/2017/06/03/a-new-cancer-drug-helped-almost-everyone-who-took-it-almost-heres-what-it-teaches-us/#2a44ee2f6b25>
2. <https://www.youtube.com/watch?v=UwbuW7oK8rk>
3. <https://www.youtube.com/watch?v=qxXRKVompI8>

1.3. Real-world/Business objectives and constraints.

- No low-latency requirement.
- Interpretability is important.
- Errors can be very costly.
- Probability of a data-point belonging to each class is needed.

2. Machine Learning Problem Formulation

2.1. Data

2.1.1. Data Overview

- Source: <https://www.kaggle.com/c/msk-redefining-cancer-treatment/data>
- We have two data files: one contains the information about the genetic mutations and the other contains the clinical evidence (text) that human experts/pathologists use to classify the genetic mutations.

- Both these data files are have a common column called ID
- Data file's information:

```
<li>
training_variants (ID , Gene, Variations, Class)
</li>
<li>
training_text (ID, Text)
</li>
```

2.1.2. Example Data Point

training_variants

ID, Gene, Variation, Class 0, FAM58A, Truncating Mutations, 1 1, CBL, W802*, 2 2, CBL, Q249E, 2 ...

training_text

ID, Text 0 | Cyclin-dependent kinases (CDKs) regulate a variety of fundamental cellular processes. CDK10 stands out as one of the last orphan CDKs for which no activating cyclin has been identified and no kinase activity revealed. Previous work has shown that CDK10 silencing increases ETS2 (v-ets erythroblastosis virus E26 oncogene homolog 2)-driven activation of the MAPK pathway, which confers tamoxifen resistance to breast cancer cells. The precise mechanisms by which CDK10 modulates ETS2 activity, and more generally the functions of CDK10, remain elusive. Here we demonstrate that CDK10 is a cyclin-dependent kinase by identifying cyclin M as an activating cyclin. Cyclin M, an orphan cyclin, is the product of FAM58A, whose mutations cause STAR syndrome, a human developmental anomaly whose features include toe syndactyly, telecanthus, and anogenital and renal malformations. We show that STAR syndrome-associated cyclin M mutants are unable to interact with CDK10. Cyclin M silencing phenocopies CDK10 silencing in increasing c-Raf and in conferring tamoxifen resistance to breast cancer cells. CDK10/cyclin M phosphorylates ETS2 in vitro, and in cells it positively controls ETS2 degradation by the proteasome. ETS2 protein levels are increased in cells derived from a STAR patient, and this increase is attributable to decreased cyclin M levels. Altogether, our results reveal an additional regulatory mechanism for ETS2, which plays key roles in cancer and development. They also shed light on the molecular mechanisms underlying STAR syndrome. Cyclin-dependent kinases (CDKs) play a pivotal role in the control of a number of fundamental cellular processes (1). The human genome contains 21 genes encoding proteins that can be considered as members of the CDK family owing to their sequence similarity with bona fide CDKs, those known to be activated by cyclins (2). Although discovered almost 20 y ago (3, 4), CDK10 remains one of the two CDKs without an identified cyclin partner. This knowledge gap has largely impeded the exploration of its biological functions. CDK10 can act as a positive cell cycle regulator in some cells (5, 6) or as a tumor suppressor in others (7, 8). CDK10 interacts with the ETS2 (v-ets erythroblastosis virus E26 oncogene homolog 2) transcription factor and inhibits its transcriptional activity through an unknown mechanism (9). CDK10 knockdown derepresses ETS2, which increases the expression of the c-Raf protein kinase, activates the MAPK pathway, and induces resistance of MCF7 cells to tamoxifen (6). ...

2.2. Mapping the real-world problem to an ML problem

2.2.1. Type of Machine Learning Problem

There are nine different classes a genetic mutation can be classified into => Multi class

2.2.2. Performance Metric

Source: <https://www.kaggle.com/c/msk-redefining-cancer-treatment#evaluation>

Metric(s): * Multi class log-loss * Confusion matrix

2.2.3. Machine Learning Objectives and Constraints

Objective: Predict the probability of each data-point belonging to each of the nine classes.

Constraints:

- Interpretability
- Class probabilities are needed.
- Penalize the errors in class probabilities => Metric is Log-loss.
- No Latency constraints.

2.3. Train, CV and Test Datasets

Split the dataset randomly into three parts train, cross validation and test with 64%,16%, 20% of data respectively

3. Exploratory Data Analysis

```
In [1]: import pandas as pd
import matplotlib.pyplot as plt
import re
import time
import warnings
import numpy as np
from nltk.corpus import stopwords
from sklearn.decomposition import TruncatedSVD
from sklearn.preprocessing import normalize
from sklearn.feature_extraction.text import CountVectorizer
from sklearn.manifold import TSNE
import seaborn as sns
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import confusion_matrix
from sklearn.metrics.classification import accuracy_score, log_loss
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.linear_model import SGDClassifier
from imblearn.over_sampling import SMOTE
from collections import Counter
from scipy.sparse import hstack
from sklearn.multiclass import OneVsRestClassifier
from sklearn.svm import SVC
# from sklearn.cross_validation import StratifiedKFold
from sklearn.model_selection import StratifiedKFold
from collections import Counter, defaultdict
from sklearn.calibration import CalibratedClassifierCV
from sklearn.naive_bayes import MultinomialNB
from sklearn.naive_bayes import GaussianNB
from sklearn.model_selection import train_test_split
from sklearn.model_selection import GridSearchCV
import math
```

```

from sklearn.metrics import normalized_mutual_info_score
from sklearn.ensemble import RandomForestClassifier
warnings.filterwarnings("ignore")

from mlxtend.classifier import StackingClassifier

from sklearn import model_selection
from sklearn.linear_model import LogisticRegression

```

3.1. Reading Data

3.1.1. Reading Gene and Variation Data

```

In [2]: data = pd.read_csv('training/training_variants')
        print('Number of data points : ', data.shape[0])
        print('Number of features : ', data.shape[1])
        print('Features : ', data.columns.values)
        data.head()

```

```

Number of data points : 3321
Number of features : 4
Features : ['ID' 'Gene' 'Variation' 'Class']

```

```

Out[2]:
   ID  Gene  Variation  Class
0   0  FAM58A  Truncating Mutations    1
1   1   CBL      W802*           2
2   2   CBL      Q249E           2
3   3   CBL      N454D           3
4   4   CBL      L399V           4

```

training/training_variants is a comma separated file containing the description of the genetic Fields are

```

<ul>
  <li><b>ID : </b>the id of the row used to link the mutation to the clinical evidence</li>
  <li><b>Gene : </b>the gene where this genetic mutation is located </li>
  <li><b>Variation : </b>the aminoacid change for this mutations </li>
  <li><b>Class :</b> 1-9 the class this genetic mutation has been classified on</li>
</ul>

```

3.1.2. Reading Text Data

```

In [3]: # note the seprator in this file
        data_text =pd.read_csv("training/training_text",sep="\\|\\|",engine="python",names=["ID"
        print('Number of data points : ', data_text.shape[0])
        print('Number of features : ', data_text.shape[1])
        print('Features : ', data_text.columns.values)
        data_text.head()

```

Number of data points : 3321
Number of features : 2
Features : ['ID' 'TEXT']

```
Out[3]:
```

	ID	TEXT
0	0	Cyclin-dependent kinases (CDKs) regulate a var...
1	1	Abstract Background Non-small cell lung canc...
2	2	Abstract Background Non-small cell lung canc...
3	3	Recent evidence has demonstrated that acquired...
4	4	Oncogenic mutations in the monomeric Casitas B...

3.1.3. Preprocessing of text

```
In [4]: # loading stop words from nltk library
stop_words = set(stopwords.words('english'))

def nlp_preprocessing(total_text, index, column):
    if type(total_text) is not int:
        string = ""
        # replace every special char with space
        total_text = re.sub('[^a-zA-Z0-9\n]', ' ', total_text)
        # replace multiple spaces with single space
        total_text = re.sub('\s+', ' ', total_text)
        # converting all the chars into lower-case.
        total_text = total_text.lower()

        for word in total_text.split():
            # if the word is a not a stop word then retain that word from the data
            if not word in stop_words:
                string += word + " "

        data_text[column][index] = string

In [5]: #text processing stage.
start_time = time.clock()
for index, row in data_text.iterrows():
    if type(row['TEXT']) is str:
        nlp_preprocessing(row['TEXT'], index, 'TEXT')
    else:
        print("there is no text description for id:",index)
print('Time took for preprocessing the text :',time.clock() - start_time, "seconds")

there is no text description for id: 1109
there is no text description for id: 1277
there is no text description for id: 1407
there is no text description for id: 1639
there is no text description for id: 2755
```

Time took for preprocessing the text : 142.199534 seconds

```
In [6]: #merging both gene_variations and text data based on ID
result = pd.merge(data, data_text,on='ID', how='left')
result.head()
```

```
Out [6]:
```

	ID	Gene	Variation	Class	\
0	0	FAM58A	Truncating Mutations	1	
1	1	CBL	W802*	2	
2	2	CBL	Q249E	2	
3	3	CBL	N454D	3	
4	4	CBL	L399V	4	

	TEXT
0	cyclin dependent kinases cdks regulate variety...
1	abstract background non small cell lung cancer...
2	abstract background non small cell lung cancer...
3	recent evidence demonstrated acquired uniparen...
4	oncogenic mutations monomeric casitas b lineag...

```
In [7]: result[result.isnull().any(axis=1)]
```

```
Out [7]:
```

	ID	Gene	Variation	Class	TEXT
1109	1109	FANCA	S1088F	1	NaN
1277	1277	ARID5B	Truncating Mutations	1	NaN
1407	1407	FGFR3	K508M	6	NaN
1639	1639	FLT1	Amplification	6	NaN
2755	2755	BRAF	G596C	7	NaN

```
In [8]: result.loc[result['TEXT'].isnull(),'TEXT'] = result['Gene'] +' '+result['Variation']
```

```
In [9]: result[result['ID']==1109]
```

```
Out [9]:
```

	ID	Gene	Variation	Class	TEXT
1109	1109	FANCA	S1088F	1	FANCA S1088F

3.1.4. Test, Train and Cross Validation Split

3.1.4.1. Splitting data into train, test and cross validation (64:20:16)

```
In [10]: y_true = result['Class'].values
result.Gene = result.Gene.str.replace('\s+', '_')
result.Variation = result.Variation.str.replace('\s+', '_')
```

```
# split the data into test and train by maintaining same distribution of output variati
X_train, test_df, y_train, y_test = train_test_split(result, y_true, stratify=y_true,
# split the train data into train and cross validation by maintaining same distributi
train_df, cv_df, y_train, y_cv = train_test_split(X_train, y_train, stratify=y_train,
```

We split the data into train, test and cross validation data sets, preserving the ratio of class distribution in the original data set

```
In [11]: print('Number of data points in train data:', train_df.shape[0])
         print('Number of data points in test data:', test_df.shape[0])
         print('Number of data points in cross validation data:', cv_df.shape[0])
```

Number of data points in train data: 2124

Number of data points in test data: 665

Number of data points in cross validation data: 532

3.1.4.2. Distribution of y_i's in Train, Test and Cross Validation datasets

```
In [12]: # it returns a dict, keys as class labels and values as the number of data points in
        train_class_distribution = train_df['Class'].value_counts().sortlevel()
        test_class_distribution = test_df['Class'].value_counts().sortlevel()
        cv_class_distribution = cv_df['Class'].value_counts().sortlevel()

        my_colors = 'rgbkymc'
        train_class_distribution.plot(kind='bar')
        plt.xlabel('Class')
        plt.ylabel('Data points per Class')
        plt.title('Distribution of y_i in train data')
        plt.grid()
        plt.show()

        # ref: argsort https://docs.scipy.org/doc/numpy/reference/generated/numpy.argsort.htm
        # -(train_class_distribution.values): the minus sign will give us in decreasing order
        sorted_yi = np.argsort(-train_class_distribution.values)
        for i in sorted_yi:
            print('Number of data points in class', i+1, ':', train_class_distribution.values[i])

        print('-'*80)
        my_colors = 'rgbkymc'
        test_class_distribution.plot(kind='bar')
        plt.xlabel('Class')
        plt.ylabel('Data points per Class')
        plt.title('Distribution of y_i in test data')
        plt.grid()
        plt.show()


        # ref: argsort https://docs.scipy.org/doc/numpy/reference/generated/numpy.argsort.htm
        # -(train_class_distribution.values): the minus sign will give us in decreasing order
        sorted_yi = np.argsort(-test_class_distribution.values)
        for i in sorted_yi:
            print('Number of data points in class', i+1, ':', test_class_distribution.values[i])
```

```

print('-'*80)
my_colors = 'rgbkymc'
cv_class_distribution.plot(kind='bar')
plt.xlabel('Class')
plt.ylabel('Data points per Class')
plt.title('Distribution of yi in cross validation data')
plt.grid()
plt.show()

# ref: argsort https://docs.scipy.org/doc/numpy/reference/generated/numpy.argsort.htm
# -(train_class_distribution.values): the minus sign will give us in decreasing order
sorted_yi = np.argsort(-train_class_distribution.values)
for i in sorted_yi:
    print('Number of data points in class', i+1, ': ', cv_class_distribution.values[i],


```




```

Number of data points in class 7 : 609 ( 28.672 %)
Number of data points in class 4 : 439 ( 20.669 %)
Number of data points in class 1 : 363 ( 17.09 %)
Number of data points in class 2 : 289 ( 13.606 %)
Number of data points in class 6 : 176 ( 8.286 %)
Number of data points in class 5 : 155 ( 7.298 %)
Number of data points in class 3 : 57 ( 2.684 %)
Number of data points in class 9 : 24 ( 1.13 %)
Number of data points in class 8 : 12 ( 0.565 %)

```

Number of data points in class 7 : 191 (28.722 %)
Number of data points in class 4 : 137 (20.602 %)
Number of data points in class 1 : 114 (17.143 %)
Number of data points in class 2 : 91 (13.684 %)
Number of data points in class 6 : 55 (8.271 %)
Number of data points in class 5 : 48 (7.218 %)
Number of data points in class 3 : 18 (2.707 %)
Number of data points in class 9 : 7 (1.053 %)
Number of data points in class 8 : 4 (0.602 %)



Number of data points in class 7 : 153 (28.759 %)
 Number of data points in class 4 : 110 (20.677 %)
 Number of data points in class 1 : 91 (17.105 %)
 Number of data points in class 2 : 72 (13.534 %)
 Number of data points in class 6 : 44 (8.271 %)
 Number of data points in class 5 : 39 (7.331 %)
 Number of data points in class 3 : 14 (2.632 %)
 Number of data points in class 9 : 6 (1.128 %)
 Number of data points in class 8 : 3 (0.564 %)

3.2 Prediction using a 'Random' Model

In a 'Random' Model, we generate the NINE class probabilities randomly such that they sum to 1.

```

In [13]: # This function plots the confusion matrices given y_i, y_i_hat.
def plot_confusion_matrix(test_y, predict_y):
    C = confusion_matrix(test_y, predict_y)
    # C = 9,9 matrix, each cell (i,j) represents number of points of class i are predicted as class j

    A = (((C.T)/(C.sum(axis=1))).T)
    #divid each element of the confusion matrix with the sum of elements in that column

    # C = [[1, 2],
  
```

```

#      [3, 4]]
# C.T = [[1, 3],
#      [2, 4]]
# C.sum(axis = 1)  axis=0 corresponds to columns and axis=1 corresponds to rows in
# C.sum(axis = 1) = [[3, 7]]
# ((C.T)/(C.sum(axis=1))) = [[1/3, 3/7]
#                          [2/3, 4/7]]

# ((C.T)/(C.sum(axis=1))).T = [[1/3, 2/3]
#                          [3/7, 4/7]]
# sum of row elements = 1

B = (C/C.sum(axis=0))
#divid each element of the confusion matrix with the sum of elements in that row
# C = [[1, 2],
#      [3, 4]]
# C.sum(axis = 0)  axis=0 corresponds to columns and axis=1 corresponds to rows in
# C.sum(axis = 0) = [[4, 6]]
# (C/C.sum(axis=0)) = [[1/4, 2/6],
#                      [3/4, 4/6]]

labels = [1,2,3,4,5,6,7,8,9]
# representing A in heatmap format
print("-"*20, "Confusion matrix", "-"*20)
plt.figure(figsize=(20,7))
sns.heatmap(C, annot=True, cmap="YlGnBu", fmt=".3f", xticklabels=labels, yticklabels=labels)
plt.xlabel('Predicted Class')
plt.ylabel('Original Class')
plt.show()

print("-"*20, "Precision matrix (Column Sum=1)", "-"*20)
plt.figure(figsize=(20,7))
sns.heatmap(B, annot=True, cmap="YlGnBu", fmt=".3f", xticklabels=labels, yticklabels=labels)
plt.xlabel('Predicted Class')
plt.ylabel('Original Class')
plt.show()

# representing B in heatmap format
print("-"*20, "Recall matrix (Row sum=1)", "-"*20)
plt.figure(figsize=(20,7))
sns.heatmap(A, annot=True, cmap="YlGnBu", fmt=".3f", xticklabels=labels, yticklabels=labels)
plt.xlabel('Predicted Class')
plt.ylabel('Original Class')
plt.show()

```

In [14]: # we need to generate 9 numbers and the sum of numbers should be 1
one solution is to generate 9 numbers and divide each of the numbers by their sum
ref: <https://stackoverflow.com/a/18662466/4084039>

```

test_data_len = test_df.shape[0]
cv_data_len = cv_df.shape[0]

# we create a output array that has exactly same size as the CV data
cv_predicted_y = np.zeros((cv_data_len,9))
for i in range(cv_data_len):
    rand_probs = np.random.rand(1,9)
    cv_predicted_y[i] = ((rand_probs/sum(sum(rand_probs))))[0])
print("Log loss on Cross Validation Data using Random Model",log_loss(y_cv,cv_predicted_y))

# Test-Set error.
#we create a output array that has exactly same as the test data
test_predicted_y = np.zeros((test_data_len,9))
for i in range(test_data_len):
    rand_probs = np.random.rand(1,9)
    test_predicted_y[i] = ((rand_probs/sum(sum(rand_probs))))[0])
print("Log loss on Test Data using Random Model",log_loss(y_test,test_predicted_y, eps=1e-15))


predicted_y =np.argmax(test_predicted_y, axis=1)
plot_confusion_matrix(y_test, predicted_y+1)

```

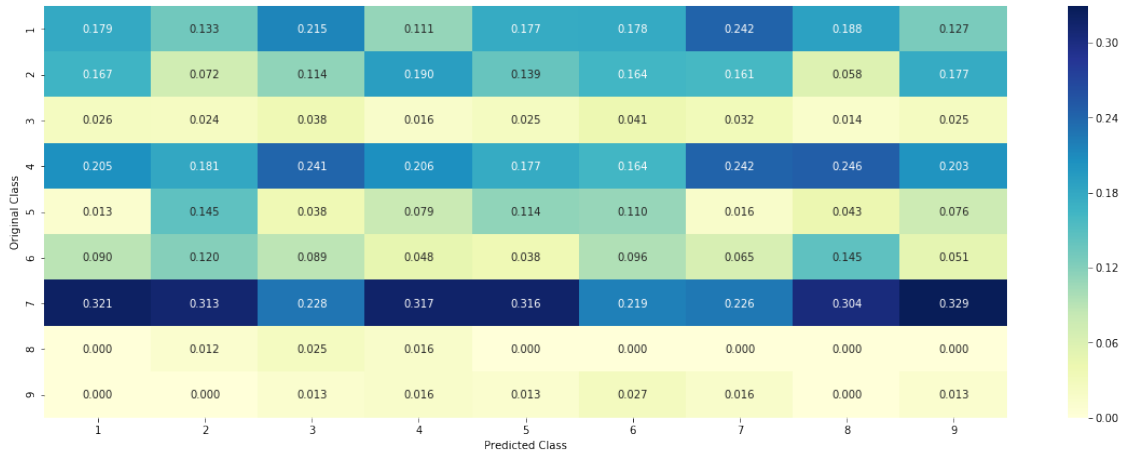
Log loss on Cross Validation Data using Random Model 2.510853375842832

Log loss on Test Data using Random Model 2.4829695422743243


----- Confusion matrix -----



----- Precision matrix (Column Sum=1) -----



----- Recall matrix (Row sum=1) -----



3.3 Univariate Analysis

```
In [15]: # code for response coding with Laplace smoothing.
# alpha : used for laplace smoothing
# feature: ['gene', 'variation']
# df: ['train_df', 'test_df', 'cv_df']
# algorithm
# -----
# Consider all unique values and the number of occurrences of given feature in train data
# build a vector (1*9), the first element = (number of times it occurred in class1 + 1)
# gv_dict is like a look up table, for every gene it stores a (1*9) representation of
# for a value of feature in df:
# if it is in train data:
```

```

# we add the vector that was stored in 'gv_dict' look up table to 'gv_fea'
# if it is not there is train:
# we add [1/9, 1/9, 1/9, 1/9,1/9, 1/9, 1/9, 1/9] to 'gv_fea'
# return 'gv_fea'
# -----

# get_gv_fea_dict: Get Gene variation Feature Dict
def get_gv_fea_dict(alpha, feature, df):
    # value_count: it contains a dict like
    # print(train_df['Gene'].value_counts())
    # output:
    #      {BRCA1      174
    #       TP53      106
    #       EGFR       86
    #       BRCA2       75
    #       PTEN       69
    #       KIT        61
    #       BRAF        60
    #       ERBB2       47
    #       PDGFRA      46
    #       ...}
    # print(train_df['Variation'].value_counts())
    # output:
    # {
    # Truncating_Mutations      63
    # Deletion                   43
    # Amplification              43
    # Fusions                    22
    # Overexpression             3
    # E17K                       3
    # Q61L                       3
    # S222D                      2
    # P130S                      2
    # ...
    # }
    value_count = train_df[feature].value_counts()

    # gv_dict : Gene Variation Dict, which contains the probability array for each gene
    gv_dict = dict()

    # denominator will contain the number of time that particular feature occurred in train
    for i, denominator in value_count.items():
        # vec will contain (p(yi==1/Gi) probability of gene/variation belongs to perturbation)
        # vec is 9 dimensional vector
        vec = []
        for k in range(1,10):
            # print(train_df.loc[(train_df['Class']==1) & (train_df['Gene']=='BRCA1')])
            #      ID      Gene      Variation      Class

```

```

# 2470 2470 BRCA1 S1715C 1
# 2486 2486 BRCA1 S1841R 1
# 2614 2614 BRCA1 M1R 1
# 2432 2432 BRCA1 L1657P 1
# 2567 2567 BRCA1 T1685A 1
# 2583 2583 BRCA1 E1660G 1
# 2634 2634 BRCA1 W1718L 1
# cls_cnt.shape[0] will return the number of rows

cls_cnt = train_df.loc[(train_df['Class']==k) & (train_df[feature]==i)]

# cls_cnt.shape[0](numerator) will contain the number of time that partic
vec.append((cls_cnt.shape[0] + alpha*10)/ (denominator + 90*alpha))

# we are adding the gene/variation to the dict as key and vec as value
gv_dict[i]=vec
return gv_dict

# Get Gene variation feature
def get_gv_feature(alpha, feature, df):
    # print(gv_dict)
    # {'BRCA1': [0.20075757575757575, 0.03787878787878788, 0.06818181818181817,
    # 'TP53': [0.32142857142857145, 0.061224489795918366, 0.061224489795918366,
    # 'EGFR': [0.056818181818181816, 0.21590909090909091, 0.0625, 0.068181818181
    # 'BRCA2': [0.13333333333333333, 0.060606060606060608, 0.060606060606060608,
    # 'PTEN': [0.069182389937106917, 0.062893081761006289, 0.069182389937106917,
    # 'KIT': [0.066225165562913912, 0.25165562913907286, 0.072847682119205295, 0
    # 'BRAF': [0.066666666666666666, 0.17999999999999999, 0.07333333333333334,
    # ...
    # }
    gv_dict = get_gv_fea_dict(alpha, feature, df)
    # value_count is similar in get_gv_fea_dict
    value_count = train_df[feature].value_counts()

    # gv_fea: Gene_variation feature, it will contain the feature for each feature va
    gv_fea = []
    # for every feature values in the given data frame we will check if it is there i
    # if not we will add [1/9,1/9,1/9,1/9,1/9,1/9,1/9,1/9,1/9] to gv_fea
    for index, row in df.iterrows():
        if row[feature] in dict(value_count).keys():
            gv_fea.append(gv_dict[row[feature]])
        else:
            gv_fea.append([1/9,1/9,1/9,1/9,1/9,1/9,1/9,1/9,1/9])
    # gv_fea.append([-1,-1,-1,-1,-1,-1,-1,-1,-1])
    return gv_fea

```

when we calculate the probability of a feature belongs to any particular class, we apply laplace smoothing

(numerator + 10*alpha) / (denominator + 90*alpha)

3.2.1 Univariate Analysis on Gene Feature

Q1. Gene, What type of feature it is ?

Ans. Gene is a categorical variable

Q2. How many categories are there and How they are distributed?

```
In [16]: unique_genes = train_df['Gene'].value_counts()
         print('Number of Unique Genes :', unique_genes.shape[0])
         # the top 10 genes that occurred most
         print(unique_genes.head(10))
```

Number of Unique Genes : 228


BRCA1	161
TP53	109
BRCA2	90
PTEN	83
EGFR	82
BRAF	62
KIT	60
ERBB2	44
ALK	40
PDGFRA	38

Name: Gene, dtype: int64


```
In [17]: print("Ans: There are", unique_genes.shape[0] , "different categories of genes in the t
```

Ans: There are 228 different categories of genes in the train data, and they are distributed as

```
In [18]: s = sum(unique_genes.values);
         h = unique_genes.values/s;
         plt.plot(h, label="Histogram of Genes")
         plt.xlabel('Index of a Gene')
         plt.ylabel('Number of Occurances')
         plt.legend()
         plt.grid()
         plt.show()
```

```
In [19]: c = np.cumsum(h)
plt.plot(c, label='Cumulative distribution of Genes')
plt.grid()
plt.legend()
plt.show()
```



Q3. How to featurize this Gene feature ?

Ans. there are two ways we can featurize this variable check out this video:
<https://www.appliedaicourse.com/course/applied-ai-course-online/lessons/handling-categorical-and-numerical-features/>

One hot Encoding

Response coding

We will choose the appropriate featurization based on the ML model we use. For this problem of multi-class classification with categorical features, one-hot encoding is better for Logistic regression while response coding is better for Random Forests.

```
In [20]: #response-coding of the Gene feature
         # alpha is used for laplace smoothing
         alpha = 1
         # train gene feature
         train_gene_feature_responseCoding = np.array(get_gv_feature(alpha, "Gene", train_df))
         # test gene feature
         test_gene_feature_responseCoding = np.array(get_gv_feature(alpha, "Gene", test_df))
         # cross validation gene feature
         cv_gene_feature_responseCoding = np.array(get_gv_feature(alpha, "Gene", cv_df))
```

```
In [21]: print("train_gene_feature_responseCoding is converted feature using response coding method")

train_gene_feature_responseCoding is converted feature using response coding method. The shape of train_gene_feature_responseCoding is (2906, 1)
```

```
In [22]: # one-hot encoding of Gene feature.
         gene_vectorizer = CountVectorizer()
         train_gene_feature_onehotCoding = gene_vectorizer.fit_transform(train_df['Gene'])
         test_gene_feature_onehotCoding = gene_vectorizer.transform(test_df['Gene'])
         cv_gene_feature_onehotCoding = gene_vectorizer.transform(cv_df['Gene'])
```

```
In [23]: train_df['Gene'].head()
```

```
Out[23]: 2906      NF2
         2966      KIT
         1675      FLT3
         2359      STK11
         2608      BRCA1
         Name: Gene, dtype: object
```

```
In [24]: gene_vectorizer.get_feature_names()
```

```
Out[24]: ['abl1',
         'acvr1',
         'ago2',
         'akt1',
         'akt2',
```

'akt3',
'alk',
'apc',
'ar',
'araf',
'arid1a',
'arid2',
'arid5b',
'asxl2',
'atm',
'atr',
'atrx',
'aurka',
'axin1',
'axl',
'b2m',
'bap1',
'bard1',
'bcl10',
'bcl2',
'bcor',
'braf',
'brca1',
'brca2',
'brip1',
'btk',
'card11',
'carm1',
'casp8',
'cbl',
'ccnd1',
'ccnd2',
'ccnd3',
'cdh1',
'cdk12',
'cdk6',
'cdkn1a',
'cdkn1b',
'cdkn2a',
'cdkn2b',
'cdkn2c',
'cebpa',
'chek2',
'cic',
'crebbp',
'ctcf',
'ctla4',
'ctnnb1',

'ddr2',
'dicer1',
'dnmt3b',
'egfr',
'elf3',
'ep300',
'epas1',
'epcam',
'erbb2',
'erbb3',
'erbb4',
'ercc2',
'ercc3',
'ercc4',
'erg',
'esr1',
'etv1',
'etv6',
'ewsr1',
'ezh2',
'fam58a',
'fanca',
'fat1',
'fbxw7',
'fgf19',
'fgf3',
'fgf4',
'fgfr1',
'fgfr2',
'fgfr3',
'fgfr4',
'flt3',
'foxa1',
'foxl2',
'foxp1',
'fubp1',
'gata3',
'gli1',
'gna11',
'gnaq',
'gnas',
'h3f3a',
'hla',
'hnf1a',
'hras',
'idh1',
'idh2',
'igf1r',

'ikbke',
'ikzf1',
'jak1',
'jak2',
'kdm5a',
'kdm5c',
'kdm6a',
'kdr',
'keap1',
'kit',
'klf4',
'kmt2a',
'kmt2b',
'kmt2c',
'kmt2d',
'knstrn',
'kras',
'lats1',
'lats2',
'map2k1',
'map2k2',
'map2k4',
'map3k1',
'mapk1',
'mdm2',
'mdm4',
'med12',
'mef2b',
'men1',
'met',
'mga',
'mlh1',
'msh2',
'msh6',
'mtor',
'myc',
'mycn',
'myd88',
'nf1',
'nf2',
'nfe2l2',
'nfkb1a',
'nkx2',
'notch1',
'notch2',
'npm1',
'nras',
'nsd1',

'ntrk1',
'ntrk2',
'ntrk3',
'nup93',
'pak1',
'pdgfra',
'pdgfrb',
'pik3ca',
'pik3cb',
'pik3cd',
'pik3r1',
'pik3r2',
'pik3r3',
'pim1',
'pms1',
'pms2',
'pole',
'ppm1d',
'ppp2r1a',
'ppp6c',
'prdm1',
'ptch1',
'pten',
'ptpn11',
'ptprd',
'ptprt',
'rab35',
'rac1',
'rad21',
'rad50',
'rad51c',
'rad51d',
'raf1',
'rara',
'rasa1',
'rb1',
'rbm10',
'ret',
'rheb',
'rhoa',
'rit1',
'rnf43',
'ros1',
'rras2',
'runx1',
'rxra',
'rybp',
'sdhb',

```

'setd2',
'sf3b1',
'shoc2',
'smad2',
'smad3',
'smad4',
'smarca4',
'smo',
'sos1',
'sox9',
'spop',
'src',
'srsf2',
'stag2',
'stat3',
'stk11',
'tert',
'tet1',
'tet2',
'tgfbr1',
'tgfbr2',
'tmprss2',
'tp53',
'tp53bp1',
'tsc1',
'tsc2',
'u2af1',
'vegfa',
'vhl',
'xrcc2',
'yap1']

```

```

In [25]: print("train_gene_feature_onehotCoding is converted feature using one-hot encoding method.")
train_gene_feature_onehotCoding is converted feature using one-hot encoding method. The shape of

```

Q4. How good is this gene feature in predicting y_i ?

There are many ways to estimate how good a feature is, in predicting y_i . One of the good methods is to build a proper ML model using just this feature. In this case, we will build a logistic regression model using only Gene feature (one hot encoded) to predict y_i .

```

In [26]: alpha = [10 ** x for x in range(-5, 1)] # hyperparam for SGD classifier.

```

```

# read more about SGDClassifier() at http://scikit-learn.org/stable/modules/generated/sklearn.linear_model.SGDClassifier.html
# -----
# default parameters
# SGDClassifier(loss=hinge, penalty=l2, alpha=0.0001, l1_ratio=0.15, fit_intercept=True,
# shuffle=True, verbose=0, epsilon=0.1, n_jobs=1, random_state=None, learning_rate=optimal)

```

```

# class_weight=None, warm_start=False, average=False, n_iter=None)

# some of methods
# fit(X, y[, coef_init, intercept_init, ])          Fit linear model with Stochastic Gradient Descent
# predict(X)          Predict class labels for samples in X.

#-----
# video link:
#-----

cv_log_error_array=[]
for i in alpha:
    clf = SGDClassifier(alpha=i, penalty='l2', loss='log', random_state=42)
    clf.fit(train_gene_feature_onehotCoding, y_train)
    sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
    sig_clf.fit(train_gene_feature_onehotCoding, y_train)
    predict_y = sig_clf.predict_proba(cv_gene_feature_onehotCoding)
    cv_log_error_array.append(log_loss(y_cv, predict_y, labels=clf.classes_, eps=1e-15))
    print('For values of alpha = ', i, "The log loss is:", log_loss(y_cv, predict_y, labels=clf.classes_, eps=1e-15))

fig, ax = plt.subplots()
ax.plot(alpha, cv_log_error_array, c='g')
for i, txt in enumerate(np.round(cv_log_error_array, 3)):
    ax.annotate((alpha[i], np.round(txt, 3)), (alpha[i], cv_log_error_array[i]))
plt.grid()
plt.title("Cross Validation Error for each alpha")
plt.xlabel("Alpha i's")
plt.ylabel("Error measure")
plt.show()

best_alpha = np.argmin(cv_log_error_array)
clf = SGDClassifier(alpha=alpha[best_alpha], penalty='l2', loss='log', random_state=42)
clf.fit(train_gene_feature_onehotCoding, y_train)
sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
sig_clf.fit(train_gene_feature_onehotCoding, y_train)

predict_y = sig_clf.predict_proba(train_gene_feature_onehotCoding)
print('For values of best alpha = ', alpha[best_alpha], "The train log loss is:", log_loss(y_train, predict_y, labels=clf.classes_, eps=1e-15))
predict_y = sig_clf.predict_proba(cv_gene_feature_onehotCoding)
print('For values of best alpha = ', alpha[best_alpha], "The cross validation log loss is:", log_loss(y_cv, predict_y, labels=clf.classes_, eps=1e-15))
predict_y = sig_clf.predict_proba(test_gene_feature_onehotCoding)
print('For values of best alpha = ', alpha[best_alpha], "The test log loss is:", log_loss(y_test, predict_y, labels=clf.classes_, eps=1e-15))


```

```

For values of alpha = 1e-05 The log loss is: 1.440470611887944
For values of alpha = 0.0001 The log loss is: 1.2680300486887892
For values of alpha = 0.001 The log loss is: 1.2830792520479086

```


For values of alpha = 0.01 The log loss is: 1.3970477449098202
 For values of alpha = 0.1 The log loss is: 1.4736833142498464
 For values of alpha = 1 The log loss is: 1.4983476782644036



For values of best alpha = 0.0001 The train log loss is: 1.02833261997304
 For values of best alpha = 0.0001 The cross validation log loss is: 1.2680300486887892
 For values of best alpha = 0.0001 The test log loss is: 1.2286393900363202

Q5. Is the Gene feature stable across all the data sets (Test, Train, Cross validation)?

Ans. Yes, it is. Otherwise, the CV and Test errors would be significantly more than train error.

In [27]: `print("Q6. How many data points in Test and CV datasets are covered by the ", unique_`

```
test_coverage=test_df[test_df['Gene'].isin(list(set(train_df['Gene'])))].shape[0]
cv_coverage=cv_df[cv_df['Gene'].isin(list(set(train_df['Gene'])))].shape[0]
```

```
print('Ans\n1. In test data',test_coverage, 'out of ',test_df.shape[0], ":",(test_coverage/test_df.shape[0])*100)
print('2. In cross validation data',cv_coverage, 'out of ',cv_df.shape[0], ":",(cv_coverage/cv_df.shape[0])*100)
```

Q6. How many data points in Test and CV datasets are covered by the 228 genes in train dataset?

Ans

1. In test data 639 out of 665 : 96.09022556390977
2. In cross validation data 507 out of 532 : 95.30075187969925

3.2.2 Univariate Analysis on Variation Feature

Q7. Variation, What type of feature is it ?

Ans. Variation is a categorical variable

Q8. How many categories are there?


```
In [28]: unique_variations = train_df['Variation'].value_counts()
         print('Number of Unique Variations :', unique_variations.shape[0])
         # the top 10 variations that occurred most
         print(unique_variations.head(10))
```

```
Number of Unique Variations : 1929
Truncating_Mutations      63
Deletion                  45
Amplification             40
Fusions                   22
Q61H                      3
T58I                      3
A146T                     2
Q209L                     2
G13V                      2
Overexpression            2
Name: Variation, dtype: int64
```

```
In [29]: print("Ans: There are", unique_variations.shape[0] ,"different categories of variation")
```


Ans: There are 1929 different categories of variations in the train data, and they are distributed as follows:

```
In [30]: s = sum(unique_variations.values);
         h = unique_variations.values/s;
         plt.plot(h, label="Histogram of Variations")
         plt.xlabel('Index of a Variation')
         plt.ylabel('Number of Occurrences')
         plt.legend()
         plt.grid()
         plt.show()
```



```
In [31]: c = np.cumsum(h)
print(c)
plt.plot(c,label='Cumulative distribution of Variations')
plt.grid()
plt.legend()
plt.show()

[0.02966102 0.05084746 0.06967985 ... 0.99905838 0.99952919 1.      ]
```



Q9. How to featurize this Variation feature ?

Ans. There are two ways we can featurize this variable check out this video:
<https://www.appliedaicourse.com/course/applied-ai-course-online/lessons/handling-categorical-and-numerical-features/>

One hot Encoding

Response coding

We will be using both these methods to featurize the Variation Feature

```
In [32]: # alpha is used for laplace smoothing
alpha = 1
# train gene feature
train_variation_feature_responseCoding = np.array(get_gv_feature(alpha, "Variation", t
# test gene feature
test_variation_feature_responseCoding = np.array(get_gv_feature(alpha, "Variation", t
# cross validation gene feature
cv_variation_feature_responseCoding = np.array(get_gv_feature(alpha, "Variation", cv_c

In [33]: print("train_variation_feature_responseCoding is a converted feature using the response
train_variation_feature_responseCoding is a converted feature using the response coding method

In [34]: # one-hot encoding of variation feature.
variation_vectorizer = CountVectorizer()
train_variation_feature_onehotCoding = variation_vectorizer.fit_transform(train_df['V
test_variation_feature_onehotCoding = variation_vectorizer.transform(test_df['Variati
cv_variation_feature_onehotCoding = variation_vectorizer.transform(cv_df['Variation'])
```

```
In [35]: print("train_variation_feature_onehotEncoded is converted feature using the onne-hot e
train_variation_feature_onehotEncoded is converted feature using the onne-hot encoding method.
```

Q10. How good is this Variation feature in predicting y_i ?
Let's build a model just like the earlier!

```
In [36]: alpha = [10 ** x for x in range(-5, 1)]

# read more about SGDClassifier() at http://scikit-learn.org/stable/modules/generated/
# -----
# default parameters
# SGDClassifier(loss=hinge, penalty=l2, alpha=0.0001, l1_ratio=0.15, fit_intercept=Tr
# shuffle=True, verbose=0, epsilon=0.1, n_jobs=1, random_state=None, learning_rate=op
# class_weight=None, warm_start=False, average=False, n_iter=None)

# some of methods
# fit(X, y[, coef_init, intercept_init, ])          Fit linear model with Stochastic Gr
# predict(X)          Predict class labels for samples in X.

#-----
# video link:
#-----

cv_log_error_array=[]
for i in alpha:
    clf = SGDClassifier(alpha=i, penalty='l2', loss='log', random_state=42)
    clf.fit(train_variation_feature_onehotCoding, y_train)

    sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
    sig_clf.fit(train_variation_feature_onehotCoding, y_train)
    predict_y = sig_clf.predict_proba(cv_variation_feature_onehotCoding)

    cv_log_error_array.append(log_loss(y_cv, predict_y, labels=clf.classes_, eps=1e-1
    print('For values of alpha = ', i, "The log loss is:",log_loss(y_cv, predict_y, l

fig, ax = plt.subplots()
ax.plot(alpha, cv_log_error_array,c='g')
for i, txt in enumerate(np.round(cv_log_error_array,3)):
    ax.annotate((alpha[i],np.round(txt,3)), (alpha[i],cv_log_error_array[i]))
plt.grid()
plt.title("Cross Validation Error for each alpha")
plt.xlabel("Alpha i's")
plt.ylabel("Error measure")
plt.show()
```


```

best_alpha = np.argmin(cv_log_error_array)
clf = SGDClassifier(alpha=alpha[best_alpha], penalty='l2', loss='log', random_state=42)
clf.fit(train_variation_feature_onehotCoding, y_train)
sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
sig_clf.fit(train_variation_feature_onehotCoding, y_train)

predict_y = sig_clf.predict_proba(train_variation_feature_onehotCoding)
print('For values of best alpha = ', alpha[best_alpha], "The train log loss is:", log_loss(y_train, predict_y))
predict_y = sig_clf.predict_proba(cv_variation_feature_onehotCoding)
print('For values of best alpha = ', alpha[best_alpha], "The cross validation log loss is:", log_loss(y_cv, predict_y))
predict_y = sig_clf.predict_proba(test_variation_feature_onehotCoding)
print('For values of best alpha = ', alpha[best_alpha], "The test log loss is:", log_loss(y_test, predict_y))

```

For values of alpha = 1e-05 The log loss is: 1.7111557675562805
 For values of alpha = 0.0001 The log loss is: 1.7007541678303986
 For values of alpha = 0.001 The log loss is: 1.7095955962408393
 For values of alpha = 0.01 The log loss is: 1.709111824934141
 For values of alpha = 0.1 The log loss is: 1.7131648413540987
 For values of alpha = 1 The log loss is: 1.7164268050341223



For values of best alpha = 0.0001 The train log loss is: 0.7358096065669432
 For values of best alpha = 0.0001 The cross validation log loss is: 1.7007541678303986

For values of best alpha = 0.0001 The test log loss is: 1.6944517261713599

Q11. Is the Variation feature stable across all the data sets (Test, Train, Cross validation)?

Ans. Not sure! But lets be very sure using the below analysis.

```
In [37]: print("Q12. How many data points are covered by total ", unique_variations.shape[0],
            test_coverage=test_df[test_df['Variation'].isin(list(set(train_df['Variation'])))]).shape[0],
            cv_coverage=cv_df[cv_df['Variation'].isin(list(set(train_df['Variation'])))]).shape[0],
            print('Ans\n1. In test data',test_coverage, 'out of',test_df.shape[0], ":",(test_coverage/test_df.shape[0]),
            print('2. In cross validation data',cv_coverage, 'out of ',cv_df.shape[0]," :", (cv_coverage/cv_df.shape[0]))
```

Q12. How many data points are covered by total 1929 genes in test and cross validation data sets?
Ans

1. In test data 77 out of 665 : 11.578947368421053
2. In cross validation data 49 out of 532 : 9.210526315789473

3.2.3 Univariate Analysis on Text Feature

1. How many unique words are present in train data?
2. How are word frequencies distributed?
3. How to featurize text field?
4. Is the text feature useful in predicting y_i?
5. Is the text feature stable across train, test and CV datasets?

```
In [38]: # cls_text is a data frame
          # for every row in data fram consider the 'TEXT'
          # split the words by space
          # make a dict with those words
          # increment its count whenever we see that word
```

```
def extract_dictionary_paddle(cls_text):
    dictionary = defaultdict(int)
    for index, row in cls_text.iterrows():
        for word in row['TEXT'].split():
            dictionary[word] +=1
    return dictionary
```

```
In [39]: import math
          #https://stackoverflow.com/a/1602964
          def get_text_responsecoding(df):
              text_feature_responseCoding = np.zeros((df.shape[0],9))
              for i in range(0,9):
                  row_index = 0
                  for index, row in df.iterrows():
                      sum_prob = 0
                      for word in row['TEXT'].split():
                          sum_prob += math.log(((dict_list[i].get(word,0)+10 )/(total_dict.get(word,0)+10)))
```

```

        text_feature_responseCoding[row_index][i] = math.exp(sum_prob/len(row['TEXT']))
        row_index += 1
    return text_feature_responseCoding

```

```

In [40]: from sklearn.feature_extraction.text import TfidfVectorizer
        from sklearn.feature_selection import chi2
        from sklearn.feature_selection import SelectKBest

        # building a CountVectorizer with all the words that occurred minimum 3 times in train
        text_vectorizer = TfidfVectorizer(min_df=3, max_features=1000)
        train_text_feature_onehotCoding = text_vectorizer.fit_transform(train_df['TEXT'])
        # getting all the feature names (words)
        train_text_features = text_vectorizer.get_feature_names()

        # train_text_feature_onehotCoding.sum(axis=0).A1 will sum every row and returns (1*number of words)
        train_text_fea_counts = train_text_feature_onehotCoding.sum(axis=0).A1

        # zip(list(text_features), text_fea_counts) will zip a word with its number of times it appears
        text_fea_dict = dict(zip(list(train_text_features), train_text_fea_counts))

        # taking TOP 1000 features
        train_text_feature_onehotCoding = SelectKBest(chi2, k=1000).fit_transform(train_text_feature_onehotCoding)

        print("Total number of unique words in train data :", len(train_text_features))
        print('train_text_feature_onehotCoding shape ==>', train_text_feature_onehotCoding.shape)

```

```

Total number of unique words in train data : 1000
train_text_feature_onehotCoding shape ==> (2124, 1000)

```

```

In [41]: train_text_features[:10], text_fea_dict.get('000')

```

```

Out[41]: (['00', '000', '01', '05', '10', '100', '11', '12', '13', '14'],
          14.783310243005301)

```

```

In [42]: dict_list = []
        # dict_list = [] contains 9 dictionaries each corresponds to a class
        for i in range(1,10):
            cls_text = train_df[train_df['Class']==i]
            # build a word dict based on the words in that class
            dict_list.append(extract_dictionary_paddle(cls_text))
            # append it to dict_list

        # dict_list[i] is build on i'th class text data
        # total_dict is build on whole training text data
        total_dict = extract_dictionary_paddle(train_df)

```



```

confuse_array = []
for i in train_text_features:
    ratios = []
    max_val = -1
    for j in range(0,9):
        ratios.append((dict_list[j][i]+10)/(total_dict[i]+90))
    confuse_array.append(ratios)
confuse_array = np.array(confuse_array)

```

```

In [43]: # dict_list[1]
# confuse_array[0],confuse_array[1],confuse_array[2],confuse_array[3],confuse_array[4]
len(confuse_array)

```

```

Out[43]: 1000

```

```

In [44]: #response coding of text features
train_text_feature_responseCoding = get_text_responsecoding(train_df)
test_text_feature_responseCoding = get_text_responsecoding(test_df)
cv_text_feature_responseCoding = get_text_responsecoding(cv_df)

```

```

In [45]: # https://stackoverflow.com/a/16202486
# we convert each row values such that they sum to 1
train_text_feature_responseCoding = (train_text_feature_responseCoding.T/train_text_f
test_text_feature_responseCoding = (test_text_feature_responseCoding.T/test_text_feat
cv_text_feature_responseCoding = (cv_text_feature_responseCoding.T/cv_text_feature_re

```

```

In [46]: # don't forget to normalize every feature
train_text_feature_onehotCoding = normalize(train_text_feature_onehotCoding, axis=0)

# we use the same vectorizer that was trained on train data
test_text_feature_onehotCoding = text_vectorizer.transform(test_df['TEXT'])
# don't forget to normalize every feature
test_text_feature_onehotCoding = normalize(test_text_feature_onehotCoding, axis=0)

# we use the same vectorizer that was trained on train data
cv_text_feature_onehotCoding = text_vectorizer.transform(cv_df['TEXT'])
# don't forget to normalize every feature
cv_text_feature_onehotCoding = normalize(cv_text_feature_onehotCoding, axis=0)

```

```

In [47]: #https://stackoverflow.com/a/2258273/4084039
sorted_text_fea_dict = dict(sorted(text_fea_dict.items(), key=lambda x: x[1] , revers
sorted_text_occur = np.array(list(sorted_text_fea_dict.values()))

```

```

In [48]: # Number of words for a given frequency.
print(Counter(sorted_text_occur))

```

```

Counter({251.44141104438117: 1, 175.341660939697: 1, 136.20164752478405: 1, 130.17587367212175

```

```

In [49]: # Train a Logistic regression+Calibration model using text features which are on-hot
alpha = [10 ** x for x in range(-5, 1)]

# read more about SGDClassifier() at http://scikit-learn.org/stable/modules/generated/
# -----
# default parameters
# SGDClassifier(loss=hinge, penalty=l2, alpha=0.0001, l1_ratio=0.15, fit_intercept=True,
# shuffle=True, verbose=0, epsilon=0.1, n_jobs=1, random_state=None, learning_rate=optimal,
# class_weight=None, warm_start=False, average=False, n_iter=None)

# some of methods
# fit(X, y[, coef_init, intercept_init, ])          Fit linear model with Stochastic Gradient Descent
# predict(X)          Predict class labels for samples in X.

#-----
# video link:
#-----

cv_log_error_array=[]
for i in alpha:
    clf = SGDClassifier(alpha=i, penalty='l2', loss='log', random_state=42)
    clf.fit(train_text_feature_onehotCoding, y_train)

    sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
    sig_clf.fit(train_text_feature_onehotCoding, y_train)
    predict_y = sig_clf.predict_proba(cv_text_feature_onehotCoding)
    cv_log_error_array.append(log_loss(y_cv, predict_y, labels=clf.classes_, eps=1e-10))
    print('For values of alpha = ', i, "The log loss is:", log_loss(y_cv, predict_y, labels=clf.classes_, eps=1e-10))

fig, ax = plt.subplots()
ax.plot(alpha, cv_log_error_array, c='g')
for i, txt in enumerate(np.round(cv_log_error_array,3)):
    ax.annotate((alpha[i],np.round(txt,3)), (alpha[i],cv_log_error_array[i]))
plt.grid()
plt.title("Cross Validation Error for each alpha")
plt.xlabel("Alpha i's")
plt.ylabel("Error measure")
plt.show()

best_alpha = np.argmin(cv_log_error_array)
clf = SGDClassifier(alpha=alpha[best_alpha], penalty='l2', loss='log', random_state=42)
clf.fit(train_text_feature_onehotCoding, y_train)
sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
sig_clf.fit(train_text_feature_onehotCoding, y_train)

predict_y = sig_clf.predict_proba(train_text_feature_onehotCoding)


```

```

print('For values of best alpha = ', alpha[best_alpha], "The train log loss is:",log_loss)
predict_y = sig_clf.predict_proba(cv_text_feature_onehotCoding)
print('For values of best alpha = ', alpha[best_alpha], "The cross validation log loss is:",log_loss)
predict_y = sig_clf.predict_proba(test_text_feature_onehotCoding)
print('For values of best alpha = ', alpha[best_alpha], "The test log loss is:",log_loss)

```

For values of alpha = 1e-05 The log loss is: 1.1990172325576594
 For values of alpha = 0.0001 The log loss is: 1.2011063515349214
 For values of alpha = 0.001 The log loss is: 1.55856814414325
 For values of alpha = 0.01 The log loss is: 2.074453907622877
 For values of alpha = 0.1 The log loss is: 2.0583897813570475
 For values of alpha = 1 The log loss is: 2.03365752438238



For values of best alpha = 1e-05 The train log loss is: 0.7781768854069671
 For values of best alpha = 1e-05 The cross validation log loss is: 1.1990172325576594
 For values of best alpha = 1e-05 The test log loss is: 1.1743373045965237

Q. Is the Text feature stable across all the data sets (Test, Train, Cross validation)?

Ans. Yes, it seems like!

```

In [51]: def get_intersec_text(df):
          df_text_vec = CountVectorizer(min_df=3)

```

```

df_text_fea = df_text_vec.fit_transform(df['TEXT'])
df_text_features = df_text_vec.get_feature_names()

df_text_fea_counts = df_text_fea.sum(axis=0).A1
df_text_fea_dict = dict(zip(list(df_text_features),df_text_fea_counts))
len1 = len(set(df_text_features))
len2 = len(set(train_text_features) & set(df_text_features))
return len1,len2

```

```

In [52]: len1,len2 = get_intersec_text(test_df)
print(np.round((len2/len1)*100, 3), "% of word of test data appeared in train data")
len1,len2 = get_intersec_text(cv_df)
print(np.round((len2/len1)*100, 3), "% of word of Cross Validation appeared in train data")

```

3.448 % of word of test data appeared in train data
3.79 % of word of Cross Validation appeared in train data

4. Machine Learning Models

```

In [52]: #Data preparation for ML models.

```

```

#Misc. functionns for ML models

```

```

def predict_and_plot_confusion_matrix(train_x, train_y, test_x, test_y, clf):
    clf.fit(train_x, train_y)
    sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
    sig_clf.fit(train_x, train_y)
    pred_y = sig_clf.predict(test_x)

    # for calculating log_loss we will provide the array of probabilities belongs to
    print("Log loss :", log_loss(test_y, sig_clf.predict_proba(test_x)))
    # calculating the number of data points that are misclassified
    print("Number of mis-classified points :", np.count_nonzero((pred_y- test_y))/test_y)
    plot_confusion_matrix(test_y, pred_y)

```

```

In [53]: def report_log_loss(train_x, train_y, test_x, test_y, clf):
    clf.fit(train_x, train_y)
    sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
    sig_clf.fit(train_x, train_y)
    sig_clf_probs = sig_clf.predict_proba(test_x)
    return log_loss(test_y, sig_clf_probs, eps=1e-15)

```

```

In [54]: # this function will be used just for naive bayes
# for the given indices, we will print the name of the features
# and we will check whether the feature present in the test point text or not
def get_impfeature_names(indices, text, gene, var, no_features):
    gene_count_vec = CountVectorizer()
    var_count_vec = CountVectorizer()

```

```

text_count_vec = TfidfVectorizer(min_df=3, max_features=1000)

gene_vec = gene_count_vec.fit(train_df['Gene'])
var_vec = var_count_vec.fit(train_df['Variation'])
text_vec = text_count_vec.fit(train_df['TEXT'])

fea1_len = len(gene_vec.get_feature_names())
fea2_len = len(var_count_vec.get_feature_names())

word_present = 0
for i,v in enumerate(indices):
    if (v < fea1_len):
        word = gene_vec.get_feature_names()[v]
        yes_no = True if word == gene else False
        if yes_no:
            word_present += 1
            print(i, "Gene feature [{}]" present in test data point [{}]" .format(w
    elif (v < fea1_len+fea2_len):
        word = var_vec.get_feature_names()[v-(fea1_len)]
        yes_no = True if word == var else False
        if yes_no:
            word_present += 1
            print(i, "variation feature [{}]" present in test data point [{}]" .form
    else:
        word = text_vec.get_feature_names()[v-(fea1_len+fea2_len)]
        yes_no = True if word in text.split() else False
        if yes_no:
            word_present += 1
            print(i, "Text feature [{}]" present in test data point [{}]" .format(w

print("Out of the top ",no_features," features ", word_present, "are present in q

```

Stacking the three types of features

In [55]: # merging gene, variance and text features

```

# building train, test and cross validation data sets
# a = [[1, 2],
#      [3, 4]]
# b = [[4, 5],
#      [6, 7]]
# hstack(a, b) = [[1, 2, 4, 5],
#                [ 3, 4, 6, 7]]

train_gene_var_onehotCoding = hstack((train_gene_feature_onehotCoding,train_variation_fea
test_gene_var_onehotCoding = hstack((test_gene_feature_onehotCoding,test_variation_fea
cv_gene_var_onehotCoding = hstack((cv_gene_feature_onehotCoding,cv_variation_feature_o

```

```

train_x_onehotCoding = hstack((train_gene_var_onehotCoding, train_text_feature_onehotCoding))
train_y = np.array(list(train_df['Class']))

test_x_onehotCoding = hstack((test_gene_var_onehotCoding, test_text_feature_onehotCoding))
test_y = np.array(list(test_df['Class']))

cv_x_onehotCoding = hstack((cv_gene_var_onehotCoding, cv_text_feature_onehotCoding))
cv_y = np.array(list(cv_df['Class']))

train_gene_var_responseCoding = np.hstack((train_gene_feature_responseCoding, train_variation_feature_responseCoding))
test_gene_var_responseCoding = np.hstack((test_gene_feature_responseCoding, test_variation_feature_responseCoding))
cv_gene_var_responseCoding = np.hstack((cv_gene_feature_responseCoding, cv_variation_feature_responseCoding))

train_x_responseCoding = np.hstack((train_gene_var_responseCoding, train_text_feature_responseCoding))
test_x_responseCoding = np.hstack((test_gene_var_responseCoding, test_text_feature_responseCoding))
cv_x_responseCoding = np.hstack((cv_gene_var_responseCoding, cv_text_feature_responseCoding))

```

```

In [56]: print("One hot encoding features :")
print("(number of data points * number of features) in train data = ", train_x_onehotCoding.shape)
print("(number of data points * number of features) in test data = ", test_x_onehotCoding.shape)
print("(number of data points * number of features) in cross validation data = ", cv_x_onehotCoding.shape)

```

```

One hot encoding features :
(number of data points * number of features) in train data = (2124, 3195)
(number of data points * number of features) in test data = (665, 3195)
(number of data points * number of features) in cross validation data = (532, 3195)

```

```

In [57]: print(" Response encoding features :")
print("(number of data points * number of features) in train data = ", train_x_responseCoding.shape)
print("(number of data points * number of features) in test data = ", test_x_responseCoding.shape)
print("(number of data points * number of features) in cross validation data = ", cv_x_responseCoding.shape)

```

```

Response encoding features :
(number of data points * number of features) in train data = (2124, 27)
(number of data points * number of features) in test data = (665, 27)
(number of data points * number of features) in cross validation data = (532, 27)

```

4.1. Base Line Model

4.1.1. Naive Bayes

4.1.1.1. Hyper parameter tuning

```

In [59]: # find more about Multinomial Naive base function here http://scikit-learn.org/stable/modules/generated/sklearn.naive\_bayes.MultinomialNB.html
# -----
# default paramters
# sklearn.naive_bayes.MultinomialNB(alpha=1.0, fit_prior=True, class_prior=None)

```

```

# some of methods of MultinomialNB()
# fit(X, y[, sample_weight])          Fit Naive Bayes classifier according to X, y
# predict(X)                          Perform classification on an array of test vectors X.
# predict_log_proba(X)                Return log-probability estimates for the test vector X.
# -----
# video link: https://www.appliedaicourse.com/course/applied-ai-course-online/lessons
# -----

# find more about CalibratedClassifierCV here at http://scikit-learn.org/stable/modules/generated/sklearn.calibration.CalibratedClassifierCV.html
# -----
# default paramters
# sklearn.calibration.CalibratedClassifierCV(base_estimator=None, method=sigmoid, cv=5)
#
# some of the methods of CalibratedClassifierCV()
# fit(X, y[, sample_weight])          Fit the calibrated model
# get_params([deep])                 Get parameters for this estimator.
# predict(X)                         Predict the target of new samples.
# predict_proba(X)                   Posterior probabilities of classification
# -----
# video link: https://www.appliedaicourse.com/course/applied-ai-course-online/lessons
# -----

alpha = [0.00001, 0.0001, 0.001, 0.1, 1, 10, 100, 1000]
cv_log_error_array = []
for i in alpha:
    print("for alpha =", i)
    clf = MultinomialNB(alpha=i)
    clf.fit(train_x_onehotCoding, train_y)
    sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
    sig_clf.fit(train_x_onehotCoding, train_y)
    sig_clf_probs = sig_clf.predict_proba(cv_x_onehotCoding)
    cv_log_error_array.append(log_loss(cv_y, sig_clf_probs, labels=clf.classes_, eps=1e-15))
    # to avoid rounding error while multiplying probabilities we use log-probability
    print("Log Loss :", log_loss(cv_y, sig_clf_probs))

print('=='*25)
fig, ax = plt.subplots()
ax.plot(np.log10(alpha), cv_log_error_array, c='g')
for i, txt in enumerate(np.round(cv_log_error_array, 3)):
    ax.annotate((alpha[i], str(txt)), (np.log10(alpha[i]), cv_log_error_array[i]))
plt.grid()
plt.xticks(np.log10(alpha))
plt.title("Cross Validation Error for each alpha")
plt.xlabel("Alpha i's")
plt.ylabel("Error measure")
plt.show()

```


```

best_alpha = np.argmin(cv_log_error_array)
clf = MultinomialNB(alpha=alpha[best_alpha])
clf.fit(train_x_onehotCoding, train_y)
sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
sig_clf.fit(train_x_onehotCoding, train_y)

predict_y = sig_clf.predict_proba(train_x_onehotCoding)
print('For values of best alpha = ', alpha[best_alpha], "The train log loss is:",log_loss(train_x_onehotCoding, predict_y))
predict_y = sig_clf.predict_proba(cv_x_onehotCoding)
print('For values of best alpha = ', alpha[best_alpha], "The cross validation log loss is:",log_loss(cv_x_onehotCoding, predict_y))
predict_y = sig_clf.predict_proba(test_x_onehotCoding)
print('For values of best alpha = ', alpha[best_alpha], "The test log loss is:",log_loss(test_x_onehotCoding, predict_y))

for alpha = 1e-05
Log Loss : 1.18618885722157
for alpha = 0.0001
Log Loss : 1.1856737416405576
for alpha = 0.001
Log Loss : 1.1837169890064192
for alpha = 0.1
Log Loss : 1.2020149288803315
for alpha = 1
Log Loss : 1.26890842893281
for alpha = 10
Log Loss : 1.4834280539786144
for alpha = 100
Log Loss : 1.4519007480265433
for alpha = 1000
Log Loss : 1.44281256571097
=====

```

For values of best alpha = 0.001 The train log loss is: 0.5234270068824562

For values of best alpha = 0.001 The cross validation log loss is: 1.1837169890064192

For values of best alpha = 0.001 The test log loss is: 1.1776485022515406

4.1.1.2. Testing the model with best hyper paramters

```
In [60]: # find more about Multinomial Naive base function here http://scikit-learn.org/stable/
# -----
# default paramters
# sklearn.naive_bayes.MultinomialNB(alpha=1.0, fit_prior=True, class_prior=None)

# some of methods of MultinomialNB()
# fit(X, y[, sample_weight])          Fit Naive Bayes classifier according to X, y
# predict(X)                          Perform classification on an array of test vectors X.
# predict_log_proba(X)                Return log-probability estimates for the test vector X.
# -----
# video link: https://www.appliedaicourse.com/course/applied-ai-course-online/lessons
# -----

# find more about CalibratedClassifierCV here at http://scikit-learn.org/stable/modules/generated/sklearn.calibration.CalibratedClassifierCV.html
# -----
# default paramters
# sklearn.calibration.CalibratedClassifierCV(base_estimator=None, method=sigmoid, cv=5)
```

```

#
# some of the methods of CalibratedClassifierCV()
# fit(X, y[, sample_weight])          Fit the calibrated model
# get_params([deep])                  Get parameters for this estimator.
# predict(X)                          Predict the target of new samples.
# predict_proba(X)                    Posterior probabilities of classification
# -----


clf = MultinomialNB(alpha=alpha[best_alpha])
clf.fit(train_x_onehotCoding, train_y)
sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
sig_clf.fit(train_x_onehotCoding, train_y)
sig_clf_probs = sig_clf.predict_proba(cv_x_onehotCoding)
# to avoid rounding error while multiplying probabilities we use log-probability estimation
print("Log Loss :", log_loss(cv_y, sig_clf_probs))
print("Number of misclassified point :", np.count_nonzero((sig_clf.predict(cv_x_onehotCoding) != cv_y)))
plot_confusion_matrix(cv_y, sig_clf.predict(cv_x_onehotCoding.toarray()))

```

Log Loss : 1.1837169890064192

Number of misclassified point : 0.37781954887218044


----- Confusion matrix -----



----- Precision matrix (Column Sum=1) -----



----- Recall matrix (Row sum=1) -----



4.1.1.3. Feature Importance, Correctly classified point

```
In [61]: test_point_index = 1
no_feature = 100
predicted_cls = sig_clf.predict(test_x_onehotCoding[test_point_index])
print("Predicted Class :", predicted_cls[0])
print("Predicted Class Probabilities:", np.round(sig_clf.predict_proba(test_x_onehotC
print("Actual Class :", test_y[test_point_index])
indices = np.argsort(-clf.coef_)[predicted_cls-1][:,no_feature]
print("-"*50)
get_impfeature_names(indices[0], test_df['TEXT'].iloc[test_point_index],test_df['Gene
```

Predicted Class : 4

Predicted Class Probabilities: [[0.2156 0.0432 0.0125 0.5692 0.0342 0.0321 0.0863 0.0029 0.004

Actual Class : 4

10 Text feature [activity] present in test data point [True]
11 Text feature [protein] present in test data point [True]
15 Text feature [function] present in test data point [True]
16 Text feature [proteins] present in test data point [True]
17 Text feature [experiments] present in test data point [True]
18 Text feature [acid] present in test data point [True]
19 Text feature [results] present in test data point [True]
20 Text feature [missense] present in test data point [True]
21 Text feature [amino] present in test data point [True]
24 Text feature [also] present in test data point [True]
25 Text feature [whereas] present in test data point [True]
26 Text feature [shown] present in test data point [True]
27 Text feature [type] present in test data point [True]
28 Text feature [whether] present in test data point [True]
29 Text feature [functional] present in test data point [True]
30 Text feature [mutations] present in test data point [True]
31 Text feature [mammalian] present in test data point [True]
33 Text feature [two] present in test data point [True]
34 Text feature [wild] present in test data point [True]
35 Text feature [indicate] present in test data point [True]
36 Text feature [reduced] present in test data point [True]
37 Text feature [previously] present in test data point [True]
38 Text feature [determined] present in test data point [True]
39 Text feature [important] present in test data point [True]
40 Text feature [described] present in test data point [True]
41 Text feature [may] present in test data point [True]
42 Text feature [30] present in test data point [True]
44 Text feature [loss] present in test data point [True]
45 Text feature [either] present in test data point [True]
46 Text feature [therefore] present in test data point [True]
48 Text feature [although] present in test data point [True]
49 Text feature [three] present in test data point [True]
50 Text feature [analysis] present in test data point [True]
51 Text feature [associated] present in test data point [True]
52 Text feature [containing] present in test data point [True]
54 Text feature [show] present in test data point [True]
55 Text feature [suppressor] present in test data point [True]
58 Text feature [determine] present in test data point [True]
59 Text feature [indicated] present in test data point [True]
61 Text feature [affect] present in test data point [True]
62 Text feature [tagged] present in test data point [True]
65 Text feature [one] present in test data point [True]
66 Text feature [suggesting] present in test data point [True]
67 Text feature [discussion] present in test data point [True]

```

68 Text feature [found] present in test data point [True]
70 Text feature [mutation] present in test data point [True]
72 Text feature [addition] present in test data point [True]
73 Text feature [however] present in test data point [True]
74 Text feature [vitro] present in test data point [True]
75 Text feature [lower] present in test data point [True]
76 Text feature [similar] present in test data point [True]
77 Text feature [using] present in test data point [True]
78 Text feature [used] present in test data point [True]
79 Text feature [analyzed] present in test data point [True]
80 Text feature [ability] present in test data point [True]
81 Text feature [generated] present in test data point [True]
82 Text feature [cells] present in test data point [True]
83 Text feature [suggest] present in test data point [True]
84 Text feature [indicates] present in test data point [True]
85 Text feature [resulting] present in test data point [True]
86 Text feature [high] present in test data point [True]
87 Text feature [10] present in test data point [True]
88 Text feature [control] present in test data point [True]
89 Text feature [contribute] present in test data point [True]
90 Text feature [buffer] present in test data point [True]
91 Text feature [effects] present in test data point [True]
92 Text feature [acids] present in test data point [True]
94 Text feature [transfection] present in test data point [True]
95 Text feature [could] present in test data point [True]
96 Text feature [bind] present in test data point [True]
97 Text feature [mutant] present in test data point [True]
98 Text feature [stability] present in test data point [True]
99 Text feature [reported] present in test data point [True]
Out of the top 100 features 73 are present in query point

```

4.1.1.4. Feature Importance, Incorrectly classified point

```

In [62]: test_point_index = 100
         no_feature = 100
         predicted_cls = sig_clf.predict(test_x_onehotCoding[test_point_index])
         print("Predicted Class :", predicted_cls[0])
         print("Predicted Class Probabilities:", np.round(sig_clf.predict_proba(test_x_onehotC
         print("Actual Class :", test_y[test_point_index])
         indices = np.argsort(-clf.coef_)[predicted_cls-1][:, :no_feature]
         print("-"*50)
         get_impfeature_names(indices[0], test_df['TEXT'].iloc[test_point_index], test_df['Gene

```

Predicted Class : 7

Predicted Class Probabilities: [[0.0579 0.0465 0.0113 0.0877 0.031 0.0293 0.7295 0.0029 0.003

Actual Class : 7

18 Text feature [activation] present in test data point [True]
19 Text feature [activated] present in test data point [True]
20 Text feature [kinase] present in test data point [True]
21 Text feature [downstream] present in test data point [True]
23 Text feature [inhibitor] present in test data point [True]
24 Text feature [cells] present in test data point [True]
25 Text feature [expressing] present in test data point [True]
26 Text feature [presence] present in test data point [True]
27 Text feature [factor] present in test data point [True]
28 Text feature [contrast] present in test data point [True]
29 Text feature [signaling] present in test data point [True]
30 Text feature [treatment] present in test data point [True]
31 Text feature [treated] present in test data point [True]
32 Text feature [independent] present in test data point [True]
33 Text feature [however] present in test data point [True]
34 Text feature [constitutive] present in test data point [True]
37 Text feature [growth] present in test data point [True]
38 Text feature [phosphorylation] present in test data point [True]
41 Text feature [10] present in test data point [True]
42 Text feature [shown] present in test data point [True]
43 Text feature [well] present in test data point [True]
44 Text feature [addition] present in test data point [True]
45 Text feature [higher] present in test data point [True]
46 Text feature [similar] present in test data point [True]
48 Text feature [compared] present in test data point [True]
49 Text feature [recently] present in test data point [True]
50 Text feature [mutations] present in test data point [True]
51 Text feature [cell] present in test data point [True]
52 Text feature [tyrosine] present in test data point [True]
53 Text feature [showed] present in test data point [True]
55 Text feature [found] present in test data point [True]
56 Text feature [may] present in test data point [True]
58 Text feature [although] present in test data point [True]
62 Text feature [previously] present in test data point [True]
64 Text feature [increased] present in test data point [True]
65 Text feature [suggest] present in test data point [True]
66 Text feature [without] present in test data point [True]
68 Text feature [pathway] present in test data point [True]
70 Text feature [survival] present in test data point [True]
71 Text feature [proliferation] present in test data point [True]
72 Text feature [pathways] present in test data point [True]
73 Text feature [mutation] present in test data point [True]
74 Text feature [20] present in test data point [True]
75 Text feature [15] present in test data point [True]
76 Text feature [enhanced] present in test data point [True]
77 Text feature [absence] present in test data point [True]
80 Text feature [including] present in test data point [True]
81 Text feature [studies] present in test data point [True]

```

83 Text feature [results] present in test data point [True]
84 Text feature [interestingly] present in test data point [True]
85 Text feature [mutant] present in test data point [True]
88 Text feature [increase] present in test data point [True]
89 Text feature [fig] present in test data point [True]
90 Text feature [either] present in test data point [True]
92 Text feature [identified] present in test data point [True]
93 Text feature [various] present in test data point [True]
94 Text feature [13] present in test data point [True]
95 Text feature [lines] present in test data point [True]
97 Text feature [described] present in test data point [True]
98 Text feature [examined] present in test data point [True]
Out of the top 100 features 60 are present in query point

```

4.2. K Nearest Neighbour Classification

4.2.1. Hyper parameter tuning

```

In [63]: # find more about KNeighborsClassifier() here http://scikit-learn.org/stable/modules/
# -----
# default parameter
# KNeighborsClassifier(n_neighbors=5, weights=uniform, algorithm=auto, leaf_size=30,
# metric=minkowski, metric_params=None, n_jobs=1, **kwargs)

# methods of
# fit(X, y) : Fit the model using X as training data and y as target values
# predict(X):Predict the class labels for the provided data
# predict_proba(X):Return probability estimates for the test data X.
#-----
# video link: https://www.appliedaicourse.com/course/applied-ai-course-online/lessons
#-----

# find more about CalibratedClassifierCV here at http://scikit-learn.org/stable/modules/
# -----
# default paramters
# sklearn.calibration.CalibratedClassifierCV(base_estimator=None, method=sigmoid, cv=
#
# some of the methods of CalibratedClassifierCV()
# fit(X, y[, sample_weight])          Fit the calibrated model
# get_params([deep])                  Get parameters for this estimator.
# predict(X)                          Predict the target of new samples.
# predict_proba(X)                    Posterior probabilities of classification
#-----
# video link:
#-----

```

```

alpha = [5, 11, 15, 21, 31, 41, 51, 99]
cv_log_error_array = []
for i in alpha:
    print("for alpha =", i)
    clf = KNeighborsClassifier(n_neighbors=i)
    clf.fit(train_x_responseCoding, train_y)
    sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
    sig_clf.fit(train_x_responseCoding, train_y)
    sig_clf_probs = sig_clf.predict_proba(cv_x_responseCoding)
    cv_log_error_array.append(log_loss(cv_y, sig_clf_probs, labels=clf.classes_, eps=
    # to avoid rounding error while multiplying probabilities we use log-probability e
    print("Log Loss :", log_loss(cv_y, sig_clf_probs))

fig, ax = plt.subplots()
ax.plot(alpha, cv_log_error_array, c='g')
for i, txt in enumerate(np.round(cv_log_error_array, 3)):
    ax.annotate((alpha[i], str(txt)), (alpha[i], cv_log_error_array[i]))
plt.grid()
plt.title("Cross Validation Error for each alpha")
plt.xlabel("Alpha i's")
plt.ylabel("Error measure")
plt.show()

best_alpha = np.argmin(cv_log_error_array)
clf = KNeighborsClassifier(n_neighbors=alpha[best_alpha])
clf.fit(train_x_responseCoding, train_y)
sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
sig_clf.fit(train_x_responseCoding, train_y)

predict_y = sig_clf.predict_proba(train_x_responseCoding)
print('For values of best alpha = ', alpha[best_alpha], "The train log loss is:", log_
predict_y = sig_clf.predict_proba(cv_x_responseCoding)
print('For values of best alpha = ', alpha[best_alpha], "The cross validation log loss
predict_y = sig_clf.predict_proba(test_x_responseCoding)
print('For values of best alpha = ', alpha[best_alpha], "The test log loss is:", log_

for alpha = 5
Log Loss : 1.0046408472971262
for alpha = 11
Log Loss : 1.0129559717309329
for alpha = 15
Log Loss : 1.01161728525124
for alpha = 21
Log Loss : 1.0275526018549699
for alpha = 31
Log Loss : 1.020394292823489
for alpha = 41


```



```

Log Loss : 1.0270641321706675
for alpha = 51
Log Loss : 1.0435689046882037
for alpha = 99
Log Loss : 1.0654120600196628

```



```

For values of best alpha = 5 The train log loss is: 0.507996008911384
For values of best alpha = 5 The cross validation log loss is: 1.0046408472971262
For values of best alpha = 5 The test log loss is: 1.0264781602437116

```

4.2.2. Testing the model with best hyper paramters

```

In [64]: # find more about KNeighborsClassifier() here http://scikit-learn.org/stable/modules/
# -----
# default parameter
# KNeighborsClassifier(n_neighbors=5, weights=uniform, algorithm=auto, leaf_size=30,
# metric=minkowski, metric_params=None, n_jobs=1, **kwargs)

# methods of
# fit(X, y) : Fit the model using X as training data and y as target values
# predict(X):Predict the class labels for the provided data
# predict_proba(X):Return probability estimates for the test data X.
#-----


```

```
# video link: https://www.appliedaicourse.com/course/applied-ai-course-online/lessons
#-----
clf = KNeighborsClassifier(n_neighbors=alpha[best_alpha])
predict_and_plot_confusion_matrix(train_x_responseCoding, train_y, cv_x_responseCoding)
```

Log loss : 1.0046408472971262

Number of mis-classified points : 0.36278195488721804


----- Confusion matrix -----



----- Precision matrix (Column Sum=1) -----



----- Recall matrix (Row sum=1) -----



4.2.3. Sample Query point -1

```
In [65]: clf = KNeighborsClassifier(n_neighbors=alpha[best_alpha])
         clf.fit(train_x_responseCoding, train_y)
         sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
         sig_clf.fit(train_x_responseCoding, train_y)

         test_point_index = 1
         predicted_cls = sig_clf.predict(test_x_responseCoding[0].reshape(1,-1))
         print("Predicted Class :", predicted_cls[0])
         print("Actual Class :", test_y[test_point_index])
         neighbors = clf.kneighbors(test_x_responseCoding[test_point_index].reshape(1, -1), al
         print("The ", alpha[best_alpha], " nearest neighbours of the test points belongs to clas
         print("Fequency of nearest points :", Counter(train_y[neighbors[1][0]]))
```

Predicted Class : 1

Actual Class : 4

The 5 nearest neighbours of the test points belongs to classes [4 4 4 4 4]

Fequency of nearest points : Counter({4: 5})

4.2.4. Sample Query Point-2

```
In [66]: clf = KNeighborsClassifier(n_neighbors=alpha[best_alpha])
         clf.fit(train_x_responseCoding, train_y)
         sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
         sig_clf.fit(train_x_responseCoding, train_y)

         test_point_index = 100

         predicted_cls = sig_clf.predict(test_x_responseCoding[test_point_index].reshape(1,-1))
         print("Predicted Class :", predicted_cls[0])
         print("Actual Class :", test_y[test_point_index])
```

```

neighbors = clf.kneighbors(test_x_responseCoding[test_point_index].reshape(1, -1), al
print("the k value for knn is",alpha[best_alpha],"and the nearest neighbours of the t
print("Fequency of nearest points :",Counter(train_y[neighbors[1][0]]))

```

Predicted Class : 7

Actual Class : 7

the k value for knn is 5 and the nearest neighbours of the test points belongs to classes [2 7

Fequency of nearest points : Counter({7: 4, 2: 1})

4.3. Logistic Regression

4.3.1. With Class balancing

4.3.1.1. Hyper paramter tuning

```

In [67]: # read more about SGDClassifier() at http://scikit-learn.org/stable/modules/generated,
# -----
# default parameters
# SGDClassifier(loss=hinge, penalty=l2, alpha=0.0001, l1_ratio=0.15, fit_intercept=Tr
# shuffle=True, verbose=0, epsilon=0.1, n_jobs=1, random_state=None, learning_rate=op
# class_weight=None, warm_start=False, average=False, n_iter=None)

# some of methods
# fit(X, y[, coef_init, intercept_init, ])          Fit linear model with Stochastic Gr
# predict(X)          Predict class labels for samples in X.

#-----
# video link: https://www.appliedaicourse.com/course/applied-ai-course-online/lessons,
#-----

# find more about CalibratedClassifierCV here at http://scikit-learn.org/stable/modul
# -----
# default paramters
# sklearn.calibration.CalibratedClassifierCV(base_estimator=None, method=sigmoid, cv=
#
# some of the methods of CalibratedClassifierCV()
# fit(X, y[, sample_weight])          Fit the calibrated model
# get_params([deep])          Get parameters for this estimator.
# predict(X)          Predict the target of new samples.
# predict_proba(X)          Posterior probabilities of classification
#-----
# video link:
#-----

alpha = [10 ** x for x in range(-6, 3)]
print('alpha ',alpha)
cv_log_error_array = []
for i in alpha:

```

```

print("for alpha =", i)
clf = SGDClassifier(class_weight='balanced', alpha=i, penalty='l2', loss='log', r
clf.fit(train_x_onehotCoding, train_y)
sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
sig_clf.fit(train_x_onehotCoding, train_y)
sig_clf_probs = sig_clf.predict_proba(cv_x_onehotCoding)
cv_log_error_array.append(log_loss(cv_y, sig_clf_probs, labels=clf.classes_, eps=
# to avoid rounding error while multiplying probabilities we use log-probability e
print("Log Loss :", log_loss(cv_y, sig_clf_probs))

fig, ax = plt.subplots()
ax.plot(alpha, cv_log_error_array, c='g')
for i, txt in enumerate(np.round(cv_log_error_array, 3)):
    ax.annotate((alpha[i], str(txt)), (alpha[i], cv_log_error_array[i]))
plt.grid()
plt.title("Cross Validation Error for each alpha")
plt.xlabel("Alpha i's")
plt.ylabel("Error measure")
plt.show()

best_alpha = np.argmin(cv_log_error_array)
clf = SGDClassifier(class_weight='balanced', alpha=alpha[best_alpha], penalty='l2', l
clf.fit(train_x_onehotCoding, train_y)
sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
sig_clf.fit(train_x_onehotCoding, train_y)

predict_y = sig_clf.predict_proba(train_x_onehotCoding)
print('For values of best alpha = ', alpha[best_alpha], "The train log loss is:", log_
predict_y = sig_clf.predict_proba(cv_x_onehotCoding)
print('For values of best alpha = ', alpha[best_alpha], "The cross validation log loss
predict_y = sig_clf.predict_proba(test_x_onehotCoding)
print('For values of best alpha = ', alpha[best_alpha], "The test log loss is:", log_

```

```


alpha [1e-06, 1e-05, 0.0001, 0.001, 0.01, 0.1, 1, 10, 100]
for alpha = 1e-06
Log Loss : 1.1226145213962426
for alpha = 1e-05
Log Loss : 1.0592834881832796
for alpha = 0.0001
Log Loss : 0.9733629323289568
for alpha = 0.001
Log Loss : 1.0058288780621785
for alpha = 0.01
Log Loss : 1.2268252056735183
for alpha = 0.1
Log Loss : 1.6287446598805144
for alpha = 1

```

```

Log Loss : 1.771079017018445
for alpha = 10
Log Loss : 1.7889595897623696
for alpha = 100
Log Loss : 1.7911262162041959

```



```

For values of best alpha = 0.0001 The train log loss is: 0.4391670289051929
For values of best alpha = 0.0001 The cross validation log loss is: 0.9733629323289568
For values of best alpha = 0.0001 The test log loss is: 0.9988417135791545

```

4.3.1.2. Testing the model with best hyper paramters

```

In [68]: # read more about SGDClassifier() at http://scikit-learn.org/stable/modules/generated,
# -----
# default parameters
# SGDClassifier(loss=hinge, penalty=l2, alpha=0.0001, l1_ratio=0.15, fit_intercept=True,
# shuffle=True, verbose=0, epsilon=0.1, n_jobs=1, random_state=None, learning_rate=optimal,
# class_weight=None, warm_start=False, average=False, n_iter=None)

# some of methods
# fit(X, y[, coef_init, intercept_init, ])          Fit linear model with Stochastic Gradient Descent
# predict(X)          Predict class labels for samples in X.


```

```
#-----
# video link: https://www.appliedaicourse.com/course/applied-ai-course-online/lessons
#-----
clf = SGDClassifier(class_weight='balanced', alpha=alpha[best_alpha], penalty='l2', l
predict_and_plot_confusion_matrix(train_x_onehotCoding, train_y, cv_x_onehotCoding, cv
```

Log loss : 0.9733629323289568

Number of mis-classified points : 0.34774436090225563


----- Confusion matrix -----



----- Precision matrix (Column Sum=1) -----



----- Recall matrix (Row sum=1) -----



4.3.1.3. Feature Importance

```
In [69]: def get_imp_feature_names(text, indices, removed_ind = []):
word_present = 0
tabulte_list = []
increasingorder_ind = 0
for i in indices:
    if i < train_gene_feature_onehotCoding.shape[1]:
        tabulte_list.append([increasingorder_ind, "Gene", "Yes"])
    elif i < 18:
        tabulte_list.append([increasingorder_ind, "Variation", "Yes"])
    if ((i > 17) & (i not in removed_ind)) :
        word = train_text_features[i]
        yes_no = True if word in text.split() else False
        if yes_no:
            word_present += 1
            tabulte_list.append([increasingorder_ind, train_text_features[i], yes_no])
        increasingorder_ind += 1
print(word_present, "most important features are present in our query point")
print("-"*50)
print("The features that are most important of the ", predicted_cls[0], " class:")
print(tabulate(tabulte_list, headers=["Index", "Feature name", "Present or Not"])).
```

4.3.1.3.1. Correctly Classified point

```
In [70]: # from tabulate import tabulate
clf = SGDClassifier(class_weight='balanced', alpha=alpha[best_alpha], penalty='l2', 1
clf.fit(train_x_onehotCoding, train_y)
test_point_index = 1
no_feature = 500
predicted_cls = sig_clf.predict(test_x_onehotCoding[test_point_index])
print("Predicted Class :", predicted_cls[0])
print("Predicted Class Probabilities:", np.round(sig_clf.predict_proba(test_x_onehotC
```



```

print("Actual Class :", test_y[test_point_index])
indices = np.argsort(-clf.coef_)[predicted_cls-1][:,no_feature]
print("-"*50)
get_impfeature_names(indices[0], test_df['TEXT'].iloc[test_point_index],test_df['Gene

```

Predicted Class : 4

Predicted Class Probabilities: [[0.2009 0.0208 0.0149 0.6981 0.0214 0.0165 0.0166 0.0046 0.0063

Actual Class : 4

```

-----
47 Text feature [mammalian] present in test data point [True]
50 Text feature [missense] present in test data point [True]
69 Text feature [tagged] present in test data point [True]
80 Text feature [suppressor] present in test data point [True]
114 Text feature [caused] present in test data point [True]
121 Text feature [tgf] present in test data point [True]
160 Text feature [show] present in test data point [True]
162 Text feature [protein] present in test data point [True]
166 Text feature [western] present in test data point [True]
171 Text feature [high] present in test data point [True]
172 Text feature [stability] present in test data point [True]
207 Text feature [age] present in test data point [True]
218 Text feature [tumorigenesis] present in test data point [True]
224 Text feature [germline] present in test data point [True]
237 Text feature [inactivation] present in test data point [True]
248 Text feature [suggesting] present in test data point [True]
256 Text feature [reduced] present in test data point [True]
271 Text feature [representative] present in test data point [True]
277 Text feature [buffer] present in test data point [True]
283 Text feature [short] present in test data point [True]
285 Text feature [represent] present in test data point [True]
290 Text feature [indicate] present in test data point [True]
306 Text feature [transfected] present in test data point [True]
311 Text feature [functional] present in test data point [True]
314 Text feature [resulting] present in test data point [True]
318 Text feature [bind] present in test data point [True]
328 Text feature [mice] present in test data point [True]
336 Text feature [due] present in test data point [True]
347 Text feature [4a] present in test data point [True]
350 Text feature [washed] present in test data point [True]
355 Text feature [localization] present in test data point [True]
360 Text feature [particular] present in test data point [True]
365 Text feature [ref] present in test data point [True]
373 Text feature [anti] present in test data point [True]
374 Text feature [often] present in test data point [True]
377 Text feature [regulation] present in test data point [True]
386 Text feature [activity] present in test data point [True]
393 Text feature [smad4] present in test data point [True]
407 Text feature [transfection] present in test data point [True]

```

```

408 Text feature [proteins] present in test data point [True]
413 Text feature [changes] present in test data point [True]
418 Text feature [loss] present in test data point [True]
427 Text feature [cases] present in test data point [True]
428 Text feature [cohort] present in test data point [True]
432 Text feature [bound] present in test data point [True]
433 Text feature [direct] present in test data point [True]
435 Text feature [genome] present in test data point [True]
448 Text feature [suggested] present in test data point [True]
450 Text feature [52] present in test data point [True]
451 Text feature [function] present in test data point [True]
455 Text feature [regions] present in test data point [True]
457 Text feature [family] present in test data point [True]
459 Text feature [characterized] present in test data point [True]
460 Text feature [comparison] present in test data point [True]
462 Text feature [lysates] present in test data point [True]
464 Text feature [27] present in test data point [True]
470 Text feature [flag] present in test data point [True]
473 Text feature [core] present in test data point [True]
476 Text feature [nuclear] present in test data point [True]
483 Text feature [deletion] present in test data point [True]
484 Text feature [despite] present in test data point [True]
488 Text feature [therefore] present in test data point [True]
491 Text feature [involved] present in test data point [True]
492 Text feature [times] present in test data point [True]
493 Text feature [recent] present in test data point [True]
494 Text feature [antibodies] present in test data point [True]
497 Text feature [44] present in test data point [True]
498 Text feature [29] present in test data point [True]
499 Text feature [altered] present in test data point [True]
Out of the top 500 features 69 are present in query point

```

4.3.1.3.2. Incorrectly Classified point

```

In [71]: test_point_index = 100
         no_feature = 500
         predicted_cls = sig_clf.predict(test_x_onehotCoding[test_point_index])
         print("Predicted Class :", predicted_cls[0])
         print("Predicted Class Probabilities:", np.round(sig_clf.predict_proba(test_x_onehotC
         print("Actual Class :", test_y[test_point_index])
         indices = np.argsort(-clf.coef_)[predicted_cls-1][:,no_feature]
         print("-"*50)
         get_impfeature_names(indices[0], test_df['TEXT'].iloc[test_point_index],test_df['Gene

Predicted Class : 7
Predicted Class Probabilities: [[0.003  0.0186 0.0022 0.0391 0.0029 0.0018 0.9298 0.0015 0.001
Actual Class : 7

```

```

-----
12 Text feature [transformed] present in test data point [True]
13 Text feature [downstream] present in test data point [True]
33 Text feature [constitutive] present in test data point [True]
37 Text feature [activated] present in test data point [True]
38 Text feature [activation] present in test data point [True]
64 Text feature [61] present in test data point [True]
65 Text feature [codon] present in test data point [True]
74 Text feature [phosphorylated] present in test data point [True]
79 Text feature [examined] present in test data point [True]
85 Text feature [factor] present in test data point [True]
94 Text feature [promote] present in test data point [True]
148 Text feature [mapk] present in test data point [True]
151 Text feature [pathways] present in test data point [True]
166 Text feature [expressing] present in test data point [True]
177 Text feature [ba] present in test data point [True]
190 Text feature [signaling] present in test data point [True]
191 Text feature [nm] present in test data point [True]
208 Text feature [cultured] present in test data point [True]
211 Text feature [f3] present in test data point [True]
223 Text feature [regulated] present in test data point [True]
225 Text feature [enhanced] present in test data point [True]
257 Text feature [presence] present in test data point [True]
294 Text feature [phosphorylation] present in test data point [True]
306 Text feature [2a] present in test data point [True]
309 Text feature [bone] present in test data point [True]
314 Text feature [phospho] present in test data point [True]
320 Text feature [inhibitor] present in test data point [True]
336 Text feature [positive] present in test data point [True]
348 Text feature [marrow] present in test data point [True]
381 Text feature [extracellular] present in test data point [True]
388 Text feature [high] present in test data point [True]
389 Text feature [treated] present in test data point [True]
391 Text feature [factors] present in test data point [True]
400 Text feature [coding] present in test data point [True]
404 Text feature [university] present in test data point [True]
420 Text feature [2001] present in test data point [True]
422 Text feature [tissue] present in test data point [True]
427 Text feature [observations] present in test data point [True]
438 Text feature [survival] present in test data point [True]
439 Text feature [leukemia] present in test data point [True]
441 Text feature [2b] present in test data point [True]
444 Text feature [cdna] present in test data point [True]
467 Text feature [interestingly] present in test data point [True]
477 Text feature [fold] present in test data point [True]
Out of the top 500 features 44 are present in query point

```

4.3.2. Without Class balancing

4.3.2.1. Hyper paramter tuning

```
In [72]: # read more about SGDClassifier() at http://scikit-learn.org/stable/modules/generated,
# -----
# default parameters
# SGDClassifier(loss=hinge, penalty=l2, alpha=0.0001, l1_ratio=0.15, fit_intercept=True,
# shuffle=True, verbose=0, epsilon=0.1, n_jobs=1, random_state=None, learning_rate=optimal,
# class_weight=None, warm_start=False, average=False, n_iter=None)

# some of methods
# fit(X, y[, coef_init, intercept_init, ])          Fit linear model with Stochastic Gradient Descent
# predict(X)          Predict class labels for samples in X.

#-----
# video link: https://www.appliedaicourse.com/course/applied-ai-course-online/lessons,
#-----

# find more about CalibratedClassifierCV here at http://scikit-learn.org/stable/modules/generated/sklearn.calibration.CalibratedClassifierCV.html
# -----
# default paramters
# sklearn.calibration.CalibratedClassifierCV(base_estimator=None, method=sigmoid, cv=5)
#
# some of the methods of CalibratedClassifierCV()
# fit(X, y[, sample_weight])          Fit the calibrated model
# get_params([deep])          Get parameters for this estimator.
# predict(X)          Predict the target of new samples.
# predict_proba(X)          Posterior probabilities of classification
#-----
# video link:
#-----

alpha = [10 ** x for x in range(-6, 1)]
cv_log_error_array = []
for i in alpha:
    print("for alpha =", i)
    clf = SGDClassifier(alpha=i, penalty='l2', loss='log', random_state=42)
    clf.fit(train_x_onehotCoding, train_y)
    sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
    sig_clf.fit(train_x_onehotCoding, train_y)
    sig_clf_probs = sig_clf.predict_proba(cv_x_onehotCoding)
    cv_log_error_array.append(log_loss(cv_y, sig_clf_probs, labels=clf.classes_, eps=1e-15))
    print("Log Loss :", log_loss(cv_y, sig_clf_probs))

fig, ax = plt.subplots()
ax.plot(alpha, cv_log_error_array, c='g')
```

```


for i, txt in enumerate(np.round(cv_log_error_array,3)):
    ax.annotate((alpha[i],str(txt)), (alpha[i],cv_log_error_array[i]))
plt.grid()
plt.title("Cross Validation Error for each alpha")
plt.xlabel("Alpha i's")
plt.ylabel("Error measure")
plt.show()

best_alpha = np.argmin(cv_log_error_array)
clf = SGDClassifier(alpha=alpha[best_alpha], penalty='l2', loss='log', random_state=42)
clf.fit(train_x_onehotCoding, train_y)
sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
sig_clf.fit(train_x_onehotCoding, train_y)

predict_y = sig_clf.predict_proba(train_x_onehotCoding)
print('For values of best alpha = ', alpha[best_alpha], "The train log loss is:",log_loss(train_y, predict_y))
predict_y = sig_clf.predict_proba(cv_x_onehotCoding)
print('For values of best alpha = ', alpha[best_alpha], "The cross validation log loss is:",log_loss(train_y, predict_y))
predict_y = sig_clf.predict_proba(test_x_onehotCoding)
print('For values of best alpha = ', alpha[best_alpha], "The test log loss is:",log_loss(test_y, predict_y))

for alpha = 1e-06
Log Loss : 1.1361780246678814
for alpha = 1e-05
Log Loss : 1.1127014922803464
for alpha = 0.0001
Log Loss : 1.019022238539661
for alpha = 0.001
Log Loss : 1.1116116790283928
for alpha = 0.01
Log Loss : 1.4654157548103004
for alpha = 0.1
Log Loss : 1.8616038845694647
for alpha = 1
Log Loss : 1.9067955301269583

```



For values of best alpha = 0.0001 The train log loss is: 0.42938512770897397
 For values of best alpha = 0.0001 The cross validation log loss is: 1.019022238539661
 For values of best alpha = 0.0001 The test log loss is: 1.0258545930382228

4.3.2.2. Testing model with best hyper parameters

```
In [73]: # read more about SGDClassifier() at http://scikit-learn.org/stable/modules/generated/
# -----
# default parameters
# SGDClassifier(loss=hinge, penalty=l2, alpha=0.0001, l1_ratio=0.15, fit_intercept=True,
# shuffle=True, verbose=0, epsilon=0.1, n_jobs=1, random_state=None, learning_rate=optimal,
# class_weight=None, warm_start=False, average=False, n_iter=None)

# some of methods
# fit(X, y[, coef_init, intercept_init, ])          Fit linear model with Stochastic Gradient Descent
# predict(X)          Predict class labels for samples in X.


#-----
# video link:
#-----

clf = SGDClassifier(alpha=alpha[best_alpha], penalty='l2', loss='log', random_state=42)
predict_and_plot_confusion_matrix(train_x_onehotCoding, train_y, cv_x_onehotCoding, cv_y)
```


Log loss : 1.019022238539661

Number of mis-classified points : 0.34962406015037595

----- Confusion matrix -----



----- Precision matrix (Column Sum=1) -----



----- Recall matrix (Row sum=1) -----



4.3.2.3. Feature Importance, Correctly Classified point

```
In [74]: clf = SGDClassifier(alpha=alpha[best_alpha], penalty='l2', loss='log', random_state=42)
clf.fit(train_x_onehotCoding, train_y)
test_point_index = 1
no_feature = 500
predicted_cls = sig_clf.predict(test_x_onehotCoding[test_point_index])
print("Predicted Class :", predicted_cls[0])
print("Predicted Class Probabilities:", np.round(sig_clf.predict_proba(test_x_onehotCoding[test_point_index]), 4))
print("Actual Class :", test_y[test_point_index])
indices = np.argsort(-clf.coef_)[predicted_cls-1][:, :no_feature]
print("-"*50)
get_impfeature_names(indices[0], test_df['TEXT'].iloc[test_point_index], test_df['Gene'])
```

Predicted Class : 4

Predicted Class Probabilities: `[[0.2046 0.0209 0.0165 0.6978 0.0205 0.016 0.0174 0.0025 0.0038]`

Actual Class : 4

```

59 Text feature [mammalian] present in test data point [True]
80 Text feature [missense] present in test data point [True]
85 Text feature [suppressor] present in test data point [True]
88 Text feature [tagged] present in test data point [True]
99 Text feature [caused] present in test data point [True]
131 Text feature [tgf] present in test data point [True]
156 Text feature [western] present in test data point [True]
166 Text feature [show] present in test data point [True]
174 Text feature [high] present in test data point [True]
177 Text feature [protein] present in test data point [True]
192 Text feature [tumorigenesis] present in test data point [True]
206 Text feature [stability] present in test data point [True]
217 Text feature [age] present in test data point [True]
246 Text feature [suggesting] present in test data point [True]

```


258 Text feature [germline] present in test data point [True]
264 Text feature [representative] present in test data point [True]
268 Text feature [represent] present in test data point [True]
270 Text feature [inactivation] present in test data point [True]
276 Text feature [reduced] present in test data point [True]
278 Text feature [bind] present in test data point [True]
302 Text feature [short] present in test data point [True]
303 Text feature [indicate] present in test data point [True]
304 Text feature [often] present in test data point [True]
307 Text feature [functional] present in test data point [True]
308 Text feature [resulting] present in test data point [True]
324 Text feature [buffer] present in test data point [True]
327 Text feature [transfected] present in test data point [True]
332 Text feature [4a] present in test data point [True]
334 Text feature [due] present in test data point [True]
336 Text feature [52] present in test data point [True]
347 Text feature [washed] present in test data point [True]
350 Text feature [anti] present in test data point [True]
352 Text feature [mice] present in test data point [True]
364 Text feature [regions] present in test data point [True]
365 Text feature [particular] present in test data point [True]
370 Text feature [regulation] present in test data point [True]
372 Text feature [localization] present in test data point [True]
376 Text feature [activity] present in test data point [True]
388 Text feature [smad4] present in test data point [True]
393 Text feature [cohort] present in test data point [True]
395 Text feature [ref] present in test data point [True]
401 Text feature [bound] present in test data point [True]
405 Text feature [cases] present in test data point [True]
406 Text feature [transfection] present in test data point [True]
410 Text feature [direct] present in test data point [True]
414 Text feature [proteins] present in test data point [True]
417 Text feature [lysates] present in test data point [True]
419 Text feature [changes] present in test data point [True]
420 Text feature [function] present in test data point [True]
423 Text feature [44] present in test data point [True]
440 Text feature [despite] present in test data point [True]
441 Text feature [therefore] present in test data point [True]
446 Text feature [genome] present in test data point [True]
447 Text feature [characterized] present in test data point [True]
448 Text feature [45] present in test data point [True]
450 Text feature [recent] present in test data point [True]
453 Text feature [suggested] present in test data point [True]
457 Text feature [27] present in test data point [True]
458 Text feature [loss] present in test data point [True]
465 Text feature [antibodies] present in test data point [True]
469 Text feature [comparison] present in test data point [True]
470 Text feature [core] present in test data point [True]

```

474 Text feature [genomic] present in test data point [True]
477 Text feature [times] present in test data point [True]
478 Text feature [29] present in test data point [True]
479 Text feature [family] present in test data point [True]
480 Text feature [determine] present in test data point [True]
481 Text feature [suggest] present in test data point [True]
482 Text feature [26] present in test data point [True]
487 Text feature [still] present in test data point [True]
488 Text feature [flag] present in test data point [True]
489 Text feature [42] present in test data point [True]
492 Text feature [considered] present in test data point [True]
493 Text feature [ha] present in test data point [True]
494 Text feature [differences] present in test data point [True]
495 Text feature [involved] present in test data point [True]
Out of the top 500 features 76 are present in query point

```

4.3.2.4. Feature Importance, Inorrectly Classified point

```

In [75]: test_point_index = 100
         no_feature = 500
         predicted_cls = sig_clf.predict(test_x_onehotCoding[test_point_index])
         print("Predicted Class :", predicted_cls[0])
         print("Predicted Class Probabilities:", np.round(sig_clf.predict_proba(test_x_onehotC
         print("Actual Class :", test_y[test_point_index])
         indices = np.argsort(-clf.coef_)[predicted_cls-1][:, :no_feature]
         print("-"*50)
         get_impfeature_names(indices[0], test_df['TEXT'].iloc[test_point_index], test_df['Gene

```

Predicted Class : 7

Predicted Class Probabilities: [[3.300e-03 1.810e-02 2.500e-03 4.550e-02 2.500e-03 1.800e-03 9
3.000e-04 2.000e-04]]

Actual Class : 7

```

-----
19 Text feature [downstream] present in test data point [True]
20 Text feature [transformed] present in test data point [True]
66 Text feature [activated] present in test data point [True]
81 Text feature [codon] present in test data point [True]
82 Text feature [examined] present in test data point [True]
83 Text feature [activation] present in test data point [True]
84 Text feature [61] present in test data point [True]
88 Text feature [phosphorylated] present in test data point [True]
96 Text feature [constitutive] present in test data point [True]
107 Text feature [factor] present in test data point [True]
142 Text feature [promote] present in test data point [True]
186 Text feature [expressing] present in test data point [True]
190 Text feature [ba] present in test data point [True]
222 Text feature [pathways] present in test data point [True]

```

```

232 Text feature [f3] present in test data point [True]
249 Text feature [enhanced] present in test data point [True]
252 Text feature [nm] present in test data point [True]
270 Text feature [signaling] present in test data point [True]
287 Text feature [2a] present in test data point [True]
288 Text feature [presence] present in test data point [True]
302 Text feature [cultured] present in test data point [True]
305 Text feature [mapk] present in test data point [True]
320 Text feature [coding] present in test data point [True]
324 Text feature [positive] present in test data point [True]
327 Text feature [regulated] present in test data point [True]
334 Text feature [phosphorylation] present in test data point [True]
359 Text feature [bone] present in test data point [True]
360 Text feature [2001] present in test data point [True]
362 Text feature [observations] present in test data point [True]
368 Text feature [high] present in test data point [True]
385 Text feature [factors] present in test data point [True]
407 Text feature [phospho] present in test data point [True]
408 Text feature [inhibitor] present in test data point [True]
414 Text feature [university] present in test data point [True]
430 Text feature [2b] present in test data point [True]
436 Text feature [extracellular] present in test data point [True]
441 Text feature [fold] present in test data point [True]
442 Text feature [exon] present in test data point [True]
450 Text feature [treated] present in test data point [True]
459 Text feature [marrow] present in test data point [True]
469 Text feature [membrane] present in test data point [True]
485 Text feature [survival] present in test data point [True]
487 Text feature [tissue] present in test data point [True]
Out of the top 500 features 43 are present in query point

```

4.4. Linear Support Vector Machines

4.4.1. Hyper parameter tuning

```

In [76]: # read more about support vector machines with linear kernels here http://scikit-learn.org

# -----
# default parameters
# SVC(C=1.0, kernel=rbf, degree=3, gamma=auto, coef0=0.0, shrinking=True, probability=False,
# cache_size=200, class_weight=None, verbose=False, max_iter=-1, decision_function_shape='raw')

# Some of methods of SVM()
# fit(X, y, [sample_weight])          Fit the SVM model according to the given training samples
# predict(X)                          Perform classification on samples in X.
# -----
# video link: https://www.appliedaicourse.com/course/applied-ai-course-online/lessons
# -----

```

```

# find more about CalibratedClassifierCV here at http://scikit-learn.org/stable/modules/generated/sklearn.calibration.CalibratedClassifierCV.html
# -----
# default paramters
# sklearn.calibration.CalibratedClassifierCV(base_estimator=None, method=sigmoid, cv=5)
#
# some of the methods of CalibratedClassifierCV()
# fit(X, y[, sample_weight])          Fit the calibrated model
# get_params([deep])                  Get parameters for this estimator.
# predict(X)                          Predict the target of new samples.
# predict_proba(X)                    Posterior probabilities of classification
#-----
# video link:
#-----

alpha = [10 ** x for x in range(-5, 3)]
cv_log_error_array = []
for i in alpha:
    print("for C =", i)
    # clf = SVC(C=i, kernel='linear', probability=True, class_weight='balanced')
    clf = SGDClassifier(class_weight='balanced', alpha=i, penalty='l2', loss='hinge')
    clf.fit(train_x_onehotCoding, train_y)
    sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
    sig_clf.fit(train_x_onehotCoding, train_y)
    sig_clf_probs = sig_clf.predict_proba(cv_x_onehotCoding)
    cv_log_error_array.append(log_loss(cv_y, sig_clf_probs, labels=clf.classes_, eps=1e-15))
    print("Log Loss :", log_loss(cv_y, sig_clf_probs))

fig, ax = plt.subplots()
ax.plot(alpha, cv_log_error_array, c='g')
for i, txt in enumerate(np.round(cv_log_error_array, 3)):
    ax.annotate((alpha[i], str(txt)), (alpha[i], cv_log_error_array[i]))
plt.grid()
plt.title("Cross Validation Error for each alpha")
plt.xlabel("Alpha i's")
plt.ylabel("Error measure")
plt.show()

best_alpha = np.argmin(cv_log_error_array)
# clf = SVC(C=i, kernel='linear', probability=True, class_weight='balanced')
clf = SGDClassifier(class_weight='balanced', alpha=alpha[best_alpha], penalty='l2', loss='hinge')
clf.fit(train_x_onehotCoding, train_y)
sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
sig_clf.fit(train_x_onehotCoding, train_y)

```

```


predict_y = sig_clf.predict_proba(train_x_onehotCoding)
print('For values of best alpha = ', alpha[best_alpha], "The train log loss is:",log_1)
predict_y = sig_clf.predict_proba(cv_x_onehotCoding)
print('For values of best alpha = ', alpha[best_alpha], "The cross validation log loss")
predict_y = sig_clf.predict_proba(test_x_onehotCoding)
print('For values of best alpha = ', alpha[best_alpha], "The test log loss is:",log_1)

```

```

for C = 1e-05
Log Loss : 1.077848175029197
for C = 0.0001
Log Loss : 1.0122200442734932
for C = 0.001
Log Loss : 1.0313244189991897
for C = 0.01
Log Loss : 1.346879225564524
for C = 0.1
Log Loss : 1.6687268229384513
for C = 1
Log Loss : 1.791770054017661
for C = 10
Log Loss : 1.7917640081207253
for C = 100
Log Loss : 1.7917640697580601

```



For values of best alpha = 0.0001 The train log loss is: 0.47877678335465024
 For values of best alpha = 0.0001 The cross validation log loss is: 1.0122200442734932
 For values of best alpha = 0.0001 The test log loss is: 1.0555794392942766

4.4.2. Testing model with best hyper parameters

```
In [77]: # read more about support vector machines with linear kernels here http://scikit-learn.org/stable/modules/linear\_svm.html

# -----
# default parameters
# SVC(C=1.0, kernel=rbf, degree=3, gamma=auto, coef0=0.0, shrinking=True, probability=True,
# cache_size=200, class_weight=None, verbose=False, max_iter=-1, decision_function_shape='raw')


# Some of methods of SVM()
# fit(X, y, [sample_weight])      Fit the SVM model according to the given training samples
# predict(X)                      Perform classification on samples in X.
# -----
# video link: https://www.appliedaicourse.com/course/applied-ai-course-online/lessons/10
# -----

# clf = SVC(C=alpha[best_alpha],kernel='linear',probability=True, class_weight='balanced')
clf = SGDClassifier(alpha=alpha[best_alpha], penalty='l2', loss='hinge', random_state=0)
predict_and_plot_confusion_matrix(train_x_onehotCoding, train_y,cv_x_onehotCoding,cv_y)
```

Log loss : 1.0122200442734932

Number of mis-classified points : 0.33646616541353386


----- Confusion matrix -----



----- Precision matrix (Column Sum=1) -----



----- Recall matrix (Row sum=1) -----



4.3.3. Feature Importance

4.3.3.1. For Correctly classified point

```
In [78]: clf = SGDClassifier(alpha=alpha[best_alpha], penalty='l2', loss='hinge', random_state=
clf.fit(train_x_onehotCoding,train_y)
test_point_index = 1
# test_point_index = 100
no_feature = 500
predicted_cls = sig_clf.predict(test_x_onehotCoding[test_point_index])
print("Predicted Class :", predicted_cls[0])
print("Predicted Class Probabilities:", np.round(sig_clf.predict_proba(test_x_onehotC
print("Actual Class :", test_y[test_point_index])
indices = np.argsort(-clf.coef_)[predicted_cls-1][:,:no_feature]
```

```
print("-"*50)
get_impfeature_names(indices[0], test_df['TEXT'].iloc[test_point_index], test_df['Gene
```

Predicted Class : 4

Predicted Class Probabilities: [[0.2326 0.0341 0.0232 0.6021 0.0275 0.0232 0.047 0.0039 0.0063

Actual Class : 4

```
-----
76 Text feature [missense] present in test data point [True]
78 Text feature [mammalian] present in test data point [True]
79 Text feature [tgf] present in test data point [True]
202 Text feature [caused] present in test data point [True]
204 Text feature [show] present in test data point [True]
205 Text feature [suppressor] present in test data point [True]
209 Text feature [tagged] present in test data point [True]
210 Text feature [suggesting] present in test data point [True]
211 Text feature [high] present in test data point [True]
213 Text feature [age] present in test data point [True]
214 Text feature [short] present in test data point [True]
217 Text feature [anti] present in test data point [True]
219 Text feature [protein] present in test data point [True]
220 Text feature [germline] present in test data point [True]
222 Text feature [due] present in test data point [True]
224 Text feature [bind] present in test data point [True]
225 Text feature [52] present in test data point [True]
229 Text feature [often] present in test data point [True]
231 Text feature [functional] present in test data point [True]
232 Text feature [stability] present in test data point [True]
233 Text feature [western] present in test data point [True]
234 Text feature [reduced] present in test data point [True]
235 Text feature [tumorigenesis] present in test data point [True]
237 Text feature [regions] present in test data point [True]
242 Text feature [45] present in test data point [True]
244 Text feature [mice] present in test data point [True]
246 Text feature [activated] present in test data point [True]
247 Text feature [3b] present in test data point [True]
249 Text feature [resulting] present in test data point [True]
251 Text feature [represent] present in test data point [True]
252 Text feature [bound] present in test data point [True]
253 Text feature [determine] present in test data point [True]
254 Text feature [cases] present in test data point [True]
255 Text feature [buffer] present in test data point [True]
257 Text feature [ref] present in test data point [True]
258 Text feature [representative] present in test data point [True]
458 Text feature [washed] present in test data point [True]
459 Text feature [indicate] present in test data point [True]
463 Text feature [4a] present in test data point [True]
466 Text feature [differences] present in test data point [True]
468 Text feature [proteins] present in test data point [True]
```



```

469 Text feature [despite] present in test data point [True]
472 Text feature [transfection] present in test data point [True]
473 Text feature [inactivation] present in test data point [True]
476 Text feature [hours] present in test data point [True]
477 Text feature [regulation] present in test data point [True]
478 Text feature [family] present in test data point [True]
481 Text feature [particular] present in test data point [True]
484 Text feature [genomic] present in test data point [True]
485 Text feature [29] present in test data point [True]
487 Text feature [mouse] present in test data point [True]
489 Text feature [smad4] present in test data point [True]
490 Text feature [12] present in test data point [True]
491 Text feature [mechanisms] present in test data point [True]
492 Text feature [amino] present in test data point [True]
493 Text feature [localization] present in test data point [True]
495 Text feature [lower] present in test data point [True]
496 Text feature [44] present in test data point [True]
497 Text feature [42] present in test data point [True]
498 Text feature [comparison] present in test data point [True]
499 Text feature [suggested] present in test data point [True]
Out of the top 500 features 61 are present in query point

```

4.3.3.2. For Incorrectly classified point

```

In [79]: test_point_index = 100
         no_feature = 500
         predicted_cls = sig_clf.predict(test_x_onehotCoding[test_point_index])
         print("Predicted Class :", predicted_cls[0])
         print("Predicted Class Probabilities:", np.round(sig_clf.predict_proba(test_x_onehotC
         print("Actual Class :", test_y[test_point_index])
         indices = np.argsort(-clf.coef_)[predicted_cls-1][:, :no_feature]
         print("-"*50)
         get_impfeature_names(indices[0], test_df['TEXT'].iloc[test_point_index], test_df['Gene

```

Predicted Class : 7

Predicted Class Probabilities: [[0.004 0.0297 0.0047 0.1028 0.0162 0.0051 0.8349 0.0011 0.0011

Actual Class : 7

```

-----
29 Text feature [transformed] present in test data point [True]
30 Text feature [downstream] present in test data point [True]
33 Text feature [phosphorylated] present in test data point [True]
232 Text feature [61] present in test data point [True]
233 Text feature [examined] present in test data point [True]
234 Text feature [activated] present in test data point [True]
239 Text feature [presence] present in test data point [True]
244 Text feature [factor] present in test data point [True]
245 Text feature [nm] present in test data point [True]

```

```

248 Text feature [expressing] present in test data point [True]
251 Text feature [ba] present in test data point [True]
253 Text feature [pathways] present in test data point [True]
257 Text feature [mapk] present in test data point [True]
258 Text feature [f3] present in test data point [True]
262 Text feature [activation] present in test data point [True]
265 Text feature [observations] present in test data point [True]
267 Text feature [2001] present in test data point [True]
268 Text feature [university] present in test data point [True]
271 Text feature [phospho] present in test data point [True]
272 Text feature [fold] present in test data point [True]
275 Text feature [positive] present in test data point [True]
277 Text feature [clinically] present in test data point [True]
278 Text feature [coding] present in test data point [True]
281 Text feature [high] present in test data point [True]
282 Text feature [01] present in test data point [True]
284 Text feature [codon] present in test data point [True]
285 Text feature [factors] present in test data point [True]
286 Text feature [lead] present in test data point [True]
288 Text feature [phosphorylation] present in test data point [True]
291 Text feature [enhanced] present in test data point [True]
292 Text feature [signaling] present in test data point [True]
293 Text feature [found] present in test data point [True]
294 Text feature [materials] present in test data point [True]
295 Text feature [regulated] present in test data point [True]
296 Text feature [phosphate] present in test data point [True]
300 Text feature [tissues] present in test data point [True]
302 Text feature [promote] present in test data point [True]
304 Text feature [14] present in test data point [True]
308 Text feature [exon] present in test data point [True]
310 Text feature [constitutive] present in test data point [True]
312 Text feature [cdna] present in test data point [True]
314 Text feature [mechanisms] present in test data point [True]
315 Text feature [wt] present in test data point [True]
318 Text feature [treated] present in test data point [True]
319 Text feature [cultured] present in test data point [True]
321 Text feature [increase] present in test data point [True]
Out of the top 500 features 46 are present in query point

```

4.5 Random Forest Classifier

4.5.1. Hyper paramter tuning (With One hot Encoding)

```

In [80]: # -----
# default parameters
# sklearn.ensemble.RandomForestClassifier(n_estimators=10, criterion=gini, max_depth=
# min_samples_leaf=1, min_weight_fraction_leaf=0.0, max_features=auto, max_leaf_nodes=
# min_impurity_split=None, bootstrap=True, oob_score=False, n_jobs=1, random_state=None

```

```

# class_weight=None)

# Some of methods of RandomForestClassifier()
# fit(X, y, [sample_weight])      Fit the SVM model according to the given training
# predict(X)                      Perform classification on samples in X.
# predict_proba (X)              Perform classification on samples in X.

# some of attributes of RandomForestClassifier()
# feature_importances_ : array of shape = [n_features]
# The feature importances (the higher, the more important the feature).

# -----
# video link: https://www.appliedaicourse.com/course/applied-ai-course-online/lessons
# -----

# find more about CalibratedClassifierCV here at http://scikit-learn.org/stable/modules/generated/sklearn.calibration.CalibratedClassifierCV.html
# -----
# default paramters
# sklearn.calibration.CalibratedClassifierCV(base_estimator=None, method=sigmoid, cv=5)
#
# some of the methods of CalibratedClassifierCV()
# fit(X, y[, sample_weight])      Fit the calibrated model
# get_params([deep])              Get parameters for this estimator.
# predict(X)                      Predict the target of new samples.
# predict_proba(X)                Posterior probabilities of classification
#-----
# video link:
#-----

alpha = [100,200,500,1000,2000]
max_depth = [5, 10]
cv_log_error_array = []
for i in alpha:
    for j in max_depth:
        print("for n_estimators = ", i,"and max depth = ", j)
        clf = RandomForestClassifier(n_estimators=i, criterion='gini', max_depth=j, random_state=0)
        clf.fit(train_x_onehotCoding, train_y)
        sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
        sig_clf.fit(train_x_onehotCoding, train_y)
        sig_clf_probs = sig_clf.predict_proba(cv_x_onehotCoding)
        cv_log_error_array.append(log_loss(cv_y, sig_clf_probs, labels=clf.classes_))
        print("Log Loss :",log_loss(cv_y, sig_clf_probs))

'''fig, ax = plt.subplots()
features = np.dot(np.array(alpha)[: ,None], np.array(max_depth)[None]).ravel()
ax.plot(features, cv_log_error_array, c='g')
for i, txt in enumerate(np.round(cv_log_error_array,3)):

```

```

        ax.annotate((alpha[int(i/2)],max_depth[int(i%2)],str(txt)), (features[i],cv_log_e
plt.grid()
plt.title("Cross Validation Error for each alpha")
plt.xlabel("Alpha i's")
plt.ylabel("Error measure")
plt.show()
'''

best_alpha = np.argmin(cv_log_error_array)
clf = RandomForestClassifier(n_estimators=alpha[int(best_alpha/2)], criterion='gini',
clf.fit(train_x_onehotCoding, train_y)
sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
sig_clf.fit(train_x_onehotCoding, train_y)

predict_y = sig_clf.predict_proba(train_x_onehotCoding)
print('For values of best estimator = ', alpha[int(best_alpha/2)], "The train log loss")
predict_y = sig_clf.predict_proba(cv_x_onehotCoding)
print('For values of best estimator = ', alpha[int(best_alpha/2)], "The cross validation")
predict_y = sig_clf.predict_proba(test_x_onehotCoding)
print('For values of best estimator = ', alpha[int(best_alpha/2)], "The test log loss")

for n_estimators = 100 and max depth = 5
Log Loss : 1.1946146355042506
for n_estimators = 100 and max depth = 10
Log Loss : 1.2546815993659284
for n_estimators = 200 and max depth = 5
Log Loss : 1.185191504867985
for n_estimators = 200 and max depth = 10
Log Loss : 1.2371062243138797
for n_estimators = 500 and max depth = 5
Log Loss : 1.1825153586877637
for n_estimators = 500 and max depth = 10
Log Loss : 1.239022076656428
for n_estimators = 1000 and max depth = 5
Log Loss : 1.1769664621965408
for n_estimators = 1000 and max depth = 10
Log Loss : 1.2323966853020778
for n_estimators = 2000 and max depth = 5
Log Loss : 1.1775684947131833
for n_estimators = 2000 and max depth = 10
Log Loss : 1.234035839976541
For values of best estimator = 1000 The train log loss is: 0.8527403405228027
For values of best estimator = 1000 The cross validation log loss is: 1.176966462196541
For values of best estimator = 1000 The test log loss is: 1.174699498033692

```

4.5.2. Testing model with best hyper parameters (One Hot Encoding)

In [81]: # -----

```

# default parameters
# sklearn.ensemble.RandomForestClassifier(n_estimators=10, criterion=gini, max_depth=
# min_samples_leaf=1, min_weight_fraction_leaf=0.0, max_features=auto, max_leaf_nodes=
# min_impurity_split=None, bootstrap=True, oob_score=False, n_jobs=1, random_state=None,
# class_weight=None)

# Some of methods of RandomForestClassifier()
# fit(X, y, [sample_weight])          Fit the SVM model according to the given training
# predict(X)                          Perform classification on samples in X.
# predict_proba (X)                  Perform classification on samples in X.

# some of attributes of RandomForestClassifier()
# feature_importances_ : array of shape = [n_features]
# The feature importances (the higher, the more important the feature).

# -----
# video link: https://www.appliedaicourse.com/course/applied-ai-course-online/lessons
# -----


clf = RandomForestClassifier(n_estimators=alpha[int(best_alpha/2)], criterion='gini',
predict_and_plot_confusion_matrix(train_x_onehotCoding, train_y,cv_x_onehotCoding,cv_y)

```


Log loss : 1.176966462196541

Number of mis-classified points : 0.40037593984962405

----- Confusion matrix -----



----- Precision matrix (Column Sum=1) -----



----- Recall matrix (Row sum=1) -----



4.5.3. Feature Importance

4.5.3.1. Correctly Classified point

```
In [82]: # test_point_index = 10
clf = RandomForestClassifier(n_estimators=alpha[int(best_alpha/2)], criterion='gini',
clf.fit(train_x_onehotCoding, train_y)
sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
sig_clf.fit(train_x_onehotCoding, train_y)

test_point_index = 1
no_feature = 100
predicted_cls = sig_clf.predict(test_x_onehotCoding[test_point_index])
print("Predicted Class :", predicted_cls[0])
```

```

print("Predicted Class Probabilities:", np.round(sig_clf.predict_proba(test_x_onehotC
print("Actual Class :", test_y[test_point_index])
indices = np.argsort(-clf.feature_importances_)
print("-"*50)
get_impfeature_names(indices[:no_feature], test_df['TEXT'].iloc[test_point_index], test

```

Predicted Class : 4

Predicted Class Probabilities: [[0.272 0.032 0.0182 0.524 0.048 0.0445 0.0468 0.0047 0.0099

Actual Class : 4

```

-----
3 Text feature [function] present in test data point [True]
4 Text feature [inhibitors] present in test data point [True]
5 Text feature [activated] present in test data point [True]
6 Text feature [suppressor] present in test data point [True]
7 Text feature [activation] present in test data point [True]
9 Text feature [phosphorylation] present in test data point [True]
10 Text feature [loss] present in test data point [True]
11 Text feature [missense] present in test data point [True]
13 Text feature [inhibitor] present in test data point [True]
17 Text feature [constitutively] present in test data point [True]
20 Text feature [receptor] present in test data point [True]
27 Text feature [protein] present in test data point [True]
29 Text feature [stability] present in test data point [True]
31 Text feature [growth] present in test data point [True]
32 Text feature [classified] present in test data point [True]
35 Text feature [functional] present in test data point [True]
38 Text feature [signaling] present in test data point [True]
39 Text feature [pathogenic] present in test data point [True]
46 Text feature [cells] present in test data point [True]
49 Text feature [cell] present in test data point [True]
54 Text feature [mammalian] present in test data point [True]
61 Text feature [expression] present in test data point [True]
62 Text feature [proteins] present in test data point [True]
63 Text feature [patients] present in test data point [True]
66 Text feature [treated] present in test data point [True]
68 Text feature [information] present in test data point [True]
69 Text feature [predicted] present in test data point [True]
72 Text feature [inactivation] present in test data point [True]
80 Text feature [advanced] present in test data point [True]
82 Text feature [activity] present in test data point [True]
86 Text feature [use] present in test data point [True]
87 Text feature [clinical] present in test data point [True]
90 Text feature [lines] present in test data point [True]
91 Text feature [factor] present in test data point [True]
93 Text feature [splice] present in test data point [True]
94 Text feature [active] present in test data point [True]
97 Text feature [dna] present in test data point [True]
Out of the top 100 features 37 are present in query point

```

4.5.3.2. Inorrectly Classified point

```
In [83]: test_point_index = 100
         no_feature = 100
         predicted_cls = sig_clf.predict(test_x_onehotCoding[test_point_index])
         print("Predicted Class :", predicted_cls[0])
         print("Predicted Class Probabilities:", np.round(sig_clf.predict_proba(test_x_onehotC
         print("Actuall Class :", test_y[test_point_index])
         indices = np.argsort(-clf.feature_importances_)
         print("-"*50)
         get_impfeature_names(indices[:no_feature], test_df['TEXT'].iloc[test_point_index], tes
```

Predicted Class : 7

Predicted Class Probabilities: [[0.0594 0.1223 0.0231 0.0655 0.0488 0.0372 0.6336 0.006 0.004

Actuall Class : 7

```
-----
0 Text feature [kinase] present in test data point [True]
2 Text feature [tyrosine] present in test data point [True]
3 Text feature [function] present in test data point [True]
5 Text feature [activated] present in test data point [True]
7 Text feature [activation] present in test data point [True]
8 Text feature [constitutive] present in test data point [True]
9 Text feature [phosphorylation] present in test data point [True]
11 Text feature [missense] present in test data point [True]
12 Text feature [treatment] present in test data point [True]
13 Text feature [inhibitor] present in test data point [True]
27 Text feature [protein] present in test data point [True]
31 Text feature [growth] present in test data point [True]
32 Text feature [classified] present in test data point [True]
35 Text feature [functional] present in test data point [True]
36 Text feature [downstream] present in test data point [True]
37 Text feature [ba] present in test data point [True]
38 Text feature [signaling] present in test data point [True]
39 Text feature [pathogenic] present in test data point [True]
40 Text feature [f3] present in test data point [True]
42 Text feature [erk] present in test data point [True]
43 Text feature [months] present in test data point [True]
44 Text feature [extracellular] present in test data point [True]
46 Text feature [cells] present in test data point [True]
49 Text feature [cell] present in test data point [True]
54 Text feature [mammalian] present in test data point [True]
61 Text feature [expression] present in test data point [True]
62 Text feature [proteins] present in test data point [True]
63 Text feature [patients] present in test data point [True]
65 Text feature [phosphatase] present in test data point [True]
66 Text feature [treated] present in test data point [True]
```



```

75 Text feature [mapk] present in test data point [True]
77 Text feature [classification] present in test data point [True]
79 Text feature [expressing] present in test data point [True]
81 Text feature [proliferation] present in test data point [True]
82 Text feature [activity] present in test data point [True]
83 Text feature [affected] present in test data point [True]
84 Text feature [phospho] present in test data point [True]
85 Text feature [ras] present in test data point [True]
86 Text feature [use] present in test data point [True]
87 Text feature [clinical] present in test data point [True]
90 Text feature [lines] present in test data point [True]
91 Text feature [factor] present in test data point [True]
94 Text feature [active] present in test data point [True]
96 Text feature [survival] present in test data point [True]
97 Text feature [dna] present in test data point [True]
Out of the top 100 features 45 are present in query point

```

4.5.3. Hyper paramter tuning (With Response Coding)

```

In [84]: # -----
# default parameters
# sklearn.ensemble.RandomForestClassifier(n_estimators=10, criterion=gini, max_depth=
# min_samples_leaf=1, min_weight_fraction_leaf=0.0, max_features=auto, max_leaf_nodes=
# min_impurity_split=None, bootstrap=True, oob_score=False, n_jobs=1, random_state=None)
# class_weight=None)

# Some of methods of RandomForestClassifier()
# fit(X, y, [sample_weight])          Fit the SVM model according to the given training
# predict(X)                          Perform classification on samples in X.
# predict_proba(X)                   Perform classification on samples in X.

# some of attributes of RandomForestClassifier()
# feature_importances_ : array of shape = [n_features]
# The feature importances (the higher, the more important the feature).

# -----
# video link: https://www.appliedaicourse.com/course/applied-ai-course-online/lessons
# -----

# find more about CalibratedClassifierCV here at http://scikit-learn.org/stable/modules/generated/sklearn.calibration.CalibratedClassifierCV.html
# -----
# default paramters
# sklearn.calibration.CalibratedClassifierCV(base_estimator=None, method=sigmoid, cv=
#
# some of the methods of CalibratedClassifierCV()
# fit(X, y[, sample_weight])          Fit the calibrated model

```

```

# get_params([deep])           Get parameters for this estimator.
# predict(X)                   Predict the target of new samples.
# predict_proba(X)             Posterior probabilities of classification
#-----
# video link:
#-----

alpha = [10,50,100,200,500,1000]
max_depth = [2,3,5,10]
cv_log_error_array = []
for i in alpha:
    for j in max_depth:
        print("for n_estimators =", i,"and max depth = ", j)
        clf = RandomForestClassifier(n_estimators=i, criterion='gini', max_depth=j, r
        clf.fit(train_x_responseCoding, train_y)
        sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
        sig_clf.fit(train_x_responseCoding, train_y)
        sig_clf_probs = sig_clf.predict_proba(cv_x_responseCoding)
        cv_log_error_array.append(log_loss(cv_y, sig_clf_probs, labels=clf.classes_, c
        print("Log Loss :",log_loss(cv_y, sig_clf_probs))
    '''

fig, ax = plt.subplots()
features = np.dot(np.array(alpha)[: ,None], np.array(max_depth)[None]).ravel()
ax.plot(features, cv_log_error_array, c='g')
for i, txt in enumerate(np.round(cv_log_error_array,3)):
    ax.annotate((alpha[int(i/4)],max_depth[int(i%4)],str(txt)), (features[i],cv_log_e
plt.grid()
plt.title("Cross Validation Error for each alpha")
plt.xlabel("Alpha i's")
plt.ylabel("Error measure")
plt.show()
'''

best_alpha = np.argmin(cv_log_error_array)
clf = RandomForestClassifier(n_estimators=alpha[int(best_alpha/4)], criterion='gini',
clf.fit(train_x_responseCoding, train_y)
sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
sig_clf.fit(train_x_responseCoding, train_y)

predict_y = sig_clf.predict_proba(train_x_responseCoding)
print('For values of best alpha = ', alpha[int(best_alpha/4)], "The train log loss is
predict_y = sig_clf.predict_proba(cv_x_responseCoding)
print('For values of best alpha = ', alpha[int(best_alpha/4)], "The cross validation
predict_y = sig_clf.predict_proba(test_x_responseCoding)
print('For values of best alpha = ', alpha[int(best_alpha/4)], "The test log loss is:

for n_estimators = 10 and max depth = 2
Log Loss : 2.1481377816860823

```

```

for n_estimators = 10 and max depth = 3
Log Loss : 1.7199970157910647
for n_estimators = 10 and max depth = 5
Log Loss : 1.5399153848549174
for n_estimators = 10 and max depth = 10
Log Loss : 1.6949864970992254
for n_estimators = 50 and max depth = 2
Log Loss : 1.6841251906601102
for n_estimators = 50 and max depth = 3
Log Loss : 1.4029288008624021
for n_estimators = 50 and max depth = 5
Log Loss : 1.2877280837142862
for n_estimators = 50 and max depth = 10
Log Loss : 1.6348110584416888
for n_estimators = 100 and max depth = 2
Log Loss : 1.563840924795056
for n_estimators = 100 and max depth = 3
Log Loss : 1.4575316493591823
for n_estimators = 100 and max depth = 5
Log Loss : 1.2865001834174978
for n_estimators = 100 and max depth = 10
Log Loss : 1.5917238299264034
for n_estimators = 200 and max depth = 2
Log Loss : 1.5945038918227343
for n_estimators = 200 and max depth = 3
Log Loss : 1.4850263331175053
for n_estimators = 200 and max depth = 5
Log Loss : 1.3033907077775777
for n_estimators = 200 and max depth = 10
Log Loss : 1.6119391844802033
for n_estimators = 500 and max depth = 2
Log Loss : 1.6567231415727905
for n_estimators = 500 and max depth = 3
Log Loss : 1.5014204746331525
for n_estimators = 500 and max depth = 5
Log Loss : 1.3193633224875094
for n_estimators = 500 and max depth = 10
Log Loss : 1.5984680369342452
for n_estimators = 1000 and max depth = 2
Log Loss : 1.6264649943982123
for n_estimators = 1000 and max depth = 3
Log Loss : 1.5212167781468002
for n_estimators = 1000 and max depth = 5
Log Loss : 1.3175892392072095
for n_estimators = 1000 and max depth = 10
Log Loss : 1.5984945627666398
For values of best alpha = 100 The train log loss is: 0.054083920108801836
For values of best alpha = 100 The cross validation log loss is: 1.286500183417498

```

For values of best alpha = 100 The test log loss is: 1.2894477844824646

4.5.4. Testing model with best hyper parameters (Response Coding)

```
In [85]: # -----
# default parameters
# sklearn.ensemble.RandomForestClassifier(n_estimators=10, criterion=gini, max_depth=
# min_samples_leaf=1, min_weight_fraction_leaf=0.0, max_features=auto, max_leaf_nodes=
# min_impurity_split=None, bootstrap=True, oob_score=False, n_jobs=1, random_state=None)
# class_weight=None)

# Some of methods of RandomForestClassifier()
# fit(X, y, [sample_weight])      Fit the SVM model according to the given training
# predict(X)                      Perform classification on samples in X.
# predict_proba(X)               Perform classification on samples in X.

# some of attributes of RandomForestClassifier()
# feature_importances_ : array of shape = [n_features]
# The feature importances (the higher, the more important the feature).

# -----
# video link: https://www.appliedaicourse.com/course/applied-ai-course-online/lessons
# -----

clf = RandomForestClassifier(max_depth=max_depth[int(best_alpha%4)], n_estimators=alpha)
predict_and_plot_confusion_matrix(train_x_responseCoding, train_y,cv_x_responseCoding)
```


Log loss : 1.2865001834164913

Number of mis-classified points : 0.49624060150375937


----- Confusion matrix -----



----- Precision matrix (Column Sum=1) -----



----- Recall matrix (Row sum=1) -----



4.5.5. Feature Importance

4.5.5.1. Correctly Classified point

```
In [86]: clf = RandomForestClassifier(n_estimators=alpha[int(best_alpha/4)], criterion='gini',
    clf.fit(train_x_responseCoding, train_y)
    sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
    sig_clf.fit(train_x_responseCoding, train_y)
```

```
test_point_index = 1
```



```
In [87]: test_point_index = 100
predicted_cls = sig_clf.predict(test_x_responseCoding[test_point_index].reshape(1,-1))
print("Predicted Class :", predicted_cls[0])
print("Predicted Class Probabilities:", np.round(sig_clf.predict_proba(test_x_responseCoding[test_point_index].reshape(1,-1)), 2))
print("Actual Class :", test_y[test_point_index])
indices = np.argsort(-clf.feature_importances_)
print("-"*50)
for i in indices:
    if i<9:
        print("Gene is important feature")
    elif i<18:
        print("Variation is important feature")
    else:
        print("Text is important feature")
```

Variation is important feature
Variation is important feature
Variation is important feature
Variation is important feature
Gene is important feature
Variation is important feature
Variation is important feature
Text is important feature
Text is important feature
Gene is important feature
Text is important feature
Text is important feature
Gene is important feature
Text is important feature
Variation is important feature
Gene is important feature
Text is important feature
Gene is important feature
Gene is important feature
Variation is important feature
Gene is important feature
Gene is important feature
Text is important feature
Text is important feature
Text is important feature
Variation is important feature
Gene is important feature

4.7.1 testing with hyper parameter tuning

```
In [88]: # read more about SGDClassifier() at http://scikit-learn.org/stable/modules/generated,
# -----
# default parameters
# SGDClassifier(loss=hinge, penalty=l2, alpha=0.0001, l1_ratio=0.15, fit_intercept=True,
# shuffle=True, verbose=0, epsilon=0.1, n_jobs=1, random_state=None, learning_rate=optimal,
# class_weight=None, warm_start=False, average=False, n_iter=None)

# some of methods
# fit(X, y[, coef_init, intercept_init, ])          Fit linear model with Stochastic Gradient Descent
# predict(X)          Predict class labels for samples in X.

#-----
# video link: https://www.appliedaicourse.com/course/applied-ai-course-online/lessons,
#-----

# read more about support vector machines with linear kernels here http://scikit-learn.org/stable/modules/generated/sklearn.svm.SVC.html
# -----
# default parameters
# SVC(C=1.0, kernel=rbf, degree=3, gamma=auto, coef0=0.0, shrinking=True, probability=True,
# cache_size=200, class_weight=None, verbose=False, max_iter=-1, decision_function_shape='raw')

# Some of methods of SVM()
# fit(X, y, [sample_weight])          Fit the SVM model according to the given training data
# predict(X)          Perform classification on samples in X.
# -----
# video link: https://www.appliedaicourse.com/course/applied-ai-course-online/lessons,
# -----

# read more about support vector machines with linear kernels here http://scikit-learn.org/stable/modules/generated/sklearn.svm.SVC.html
# -----
# default parameters
# sklearn.ensemble.RandomForestClassifier(n_estimators=10, criterion=gini, max_depth=None,
# min_samples_leaf=1, min_weight_fraction_leaf=0.0, max_features='auto', max_leaf_nodes=None,
# min_impurity_split=None, bootstrap=True, oob_score=False, n_jobs=1, random_state=None,
# class_weight=None)

# Some of methods of RandomForestClassifier()
# fit(X, y, [sample_weight])          Fit the SVM model according to the given training data
# predict(X)          Perform classification on samples in X.
# predict_proba(X)          Perform classification on samples in X.

# some of attributes of RandomForestClassifier()
# feature_importances_ : array of shape = [n_features]
# The feature importances (the higher, the more important the feature).
```



```
# -----
# video link: https://www.appliedaicourse.com/course/applied-ai-course-online/lessons
# -----
```

```
clf1 = SGDClassifier(alpha=0.001, penalty='l2', loss='log', class_weight='balanced', random_state=42)
clf1.fit(train_x_onehotCoding, train_y)
sig_clf1 = CalibratedClassifierCV(clf1, method="sigmoid")
```

```
clf2 = SGDClassifier(alpha=1, penalty='l2', loss='hinge', class_weight='balanced', random_state=42)
clf2.fit(train_x_onehotCoding, train_y)
sig_clf2 = CalibratedClassifierCV(clf2, method="sigmoid")
```

```
clf3 = MultinomialNB(alpha=0.001)
clf3.fit(train_x_onehotCoding, train_y)
sig_clf3 = CalibratedClassifierCV(clf3, method="sigmoid")
```

```
sig_clf1.fit(train_x_onehotCoding, train_y)
print("Logistic Regression : Log Loss: %0.2f" % (log_loss(cv_y, sig_clf1.predict_proba(cv_x_onehotCoding))))
sig_clf2.fit(train_x_onehotCoding, train_y)
print("Support vector machines : Log Loss: %0.2f" % (log_loss(cv_y, sig_clf2.predict_proba(cv_x_onehotCoding))))
sig_clf3.fit(train_x_onehotCoding, train_y)
print("Naive Bayes : Log Loss: %0.2f" % (log_loss(cv_y, sig_clf3.predict_proba(cv_x_onehotCoding))))
print("-"*50)
alpha = [0.0001,0.001,0.01,0.1,1,10]
best_alpha = 999
for i in alpha:
    lr = LogisticRegression(C=i)
    sclf = StackingClassifier(classifiers=[sig_clf1, sig_clf2, sig_clf3], meta_classifier=lr)
    sclf.fit(train_x_onehotCoding, train_y)
    print("Stacking Classifier : for the value of alpha: %f Log Loss: %0.3f" % (i, log_loss(cv_y, sclf.predict_proba(cv_x_onehotCoding))))
    log_error = log_loss(cv_y, sclf.predict_proba(cv_x_onehotCoding))
    if best_alpha > log_error:
        best_alpha = log_error
```

```
Logistic Regression : Log Loss: 1.01
Support vector machines : Log Loss: 1.79
Naive Bayes : Log Loss: 1.18
```

```
-----
Stacking Classifier : for the value of alpha: 0.000100 Log Loss: 2.177
Stacking Classifier : for the value of alpha: 0.001000 Log Loss: 2.031
Stacking Classifier : for the value of alpha: 0.010000 Log Loss: 1.492
Stacking Classifier : for the value of alpha: 0.100000 Log Loss: 1.156
Stacking Classifier : for the value of alpha: 1.000000 Log Loss: 1.392
Stacking Classifier : for the value of alpha: 10.000000 Log Loss: 1.893
```

4.7.2 testing the model with the best hyper parameters

```
In [89]: lr = LogisticRegression(C=0.1)
        scf = StackingClassifier(classifiers=[sig_clf1, sig_clf2, sig_clf3], meta_classifier=lr)
        scf.fit(train_x_onehotCoding, train_y)

        log_error = log_loss(train_y, scf.predict_proba(train_x_onehotCoding))
        print("Log loss (train) on the stacking classifier :",log_error)

        log_error = log_loss(cv_y, scf.predict_proba(cv_x_onehotCoding))
        print("Log loss (CV) on the stacking classifier :",log_error)

        log_error = log_loss(test_y, scf.predict_proba(test_x_onehotCoding))
        print("Log loss (test) on the stacking classifier :",log_error)

        print("Number of missclassified point :", np.count_nonzero((scf.predict(test_x_onehotCoding) != test_y)))
        plot_confusion_matrix(test_y=test_y, predict_y=scf.predict(test_x_onehotCoding))
```

Log loss (train) on the stacking classifier : 0.535551969273216

Log loss (CV) on the stacking classifier : 1.1562497109759189


Log loss (test) on the stacking classifier : 1.1719959433218796

Number of missclassified point : 0.3699248120300752


----- Confusion matrix -----



----- Precision matrix (Column Sum=1) -----



----- Recall matrix (Row sum=1) -----



4.7.3 Maximum Voting classifier

```
In [90]: #Refer:http://scikit-learn.org/stable/modules/generated/sklearn.ensemble.VotingClassifier
from sklearn.ensemble import VotingClassifier
vclf = VotingClassifier(estimators=[('lr', sig_clf1), ('svc', sig_clf2), ('rf', sig_clf3)])
vclf.fit(train_x_onehotCoding, train_y)
print("Log loss (train) on the VotingClassifier :", log_loss(train_y, vclf.predict_proba(train_x_onehotCoding)))
print("Log loss (CV) on the VotingClassifier :", log_loss(cv_y, vclf.predict_proba(cv_x_onehotCoding)))
print("Log loss (test) on the VotingClassifier :", log_loss(test_y, vclf.predict_proba(test_x_onehotCoding)))
print("Number of missclassified point :", np.count_nonzero(vclf.predict(test_x_onehotCoding) != test_y))
plot_confusion_matrix(test_y=test_y, predict_y=vclf.predict(test_x_onehotCoding))
```


Log loss (train) on the VotingClassifier : 0.8311581097949535

Log loss (CV) on the VotingClassifier : 1.1739972098170757


Log loss (test) on the VotingClassifier : 1.182575428318276

Number of missclassified point : 0.3699248120300752


----- Confusion matrix -----



----- Precision matrix (Column Sum=1) -----



----- Recall matrix (Row sum=1) -----



0.1 Bigram Count Vectorizer with Logistic Regression

```
In [91]: # building a CountVectorizer with all the words that occurred minimum 3 times in train
text_vectorizer = CountVectorizer(min_df=3,ngram_range=(1,2))
train_text_feature_onehotCoding = text_vectorizer.fit_transform(train_df['TEXT'])
# print(train_text_feature_onehotCoding.shape)
# getting all the feature names (words)
train_text_features= text_vectorizer.get_feature_names()

# train_text_feature_onehotCoding.sum(axis=0).A1 will sum every row and returns (1*nu
train_text_fea_counts = train_text_feature_onehotCoding.sum(axis=0).A1

# zip(list(text_features),text_fea_counts) will zip a word with its number of times i
text_fea_dict = dict(zip(list(train_text_features),train_text_fea_counts))

print("Total number of unique words in train data :", len(train_text_features))
```

Total number of unique words in train data : 803960

```
In [92]: train_text_features[:10], text_fea_dict.get('000')
```

```
Out[92]: (['00',
            '00 00',
            '00 0005',
            '00 006',
            '00 01',
            '00 02',
            '00 03',
            '00 05',
            '00 06',
```

```
'00 08'],  
5836)
```

```
In [93]: dict_list = []  
# dict_list=[] contains 9 dictionaries each corresponds to a class  
for i in range(1,10):  
    cls_text = train_df[train_df['Class']==i]  
    # build a word dict based on the words in that class  
    dict_list.append(extract_dictionary_paddle(cls_text))  
    # append it to dict_list  
  
# dict_list[i] is build on i'th class text data  
# total_dict is build on whole training text data  
total_dict = extract_dictionary_paddle(train_df)
```

```
confuse_array = []  
for i in train_text_features:  
    ratios = []  
    max_val = -1  
    for j in range(0,9):  
        ratios.append((dict_list[j][i]+10)/(total_dict[i]+90))  
    confuse_array.append(ratios)  
confuse_array = np.array(confuse_array)
```

```
In [94]: len(confuse_array)
```

```
Out[94]: 803960
```

```
In [95]: # don't forget to normalize every feature  
train_text_feature_onehotCoding = normalize(train_text_feature_onehotCoding, axis=0)  
  
# we use the same vectorizer that was trained on train data  
test_text_feature_onehotCoding = text_vectorizer.transform(test_df['TEXT'])  
# don't forget to normalize every feature  
test_text_feature_onehotCoding = normalize(test_text_feature_onehotCoding, axis=0)  
  
# we use the same vectorizer that was trained on train data  
cv_text_feature_onehotCoding = text_vectorizer.transform(cv_df['TEXT'])  
# don't forget to normalize every feature  
cv_text_feature_onehotCoding = normalize(cv_text_feature_onehotCoding, axis=0)
```

```
In [96]: # merging gene, variance and text features
```

```
# building train, test and cross validation data sets  
# a = [[1, 2],  
#      [3, 4]]  
# b = [[4, 5],  
#      [6, 7]]
```

```

# hstack(a, b) = [[1, 2, 4, 5],
#                [ 3, 4, 6, 7]]

train_gene_var_onehotCoding = hstack((train_gene_feature_onehotCoding, train_variation_f
test_gene_var_onehotCoding = hstack((test_gene_feature_onehotCoding, test_variation_fe
cv_gene_var_onehotCoding = hstack((cv_gene_feature_onehotCoding, cv_variation_feature_o

train_x_onehotCoding = hstack((train_gene_var_onehotCoding, train_text_feature_onehotC
train_y = np.array(list(train_df['Class']))

test_x_onehotCoding = hstack((test_gene_var_onehotCoding, test_text_feature_onehotCod
test_y = np.array(list(test_df['Class']))

cv_x_onehotCoding = hstack((cv_gene_var_onehotCoding, cv_text_feature_onehotCoding)).
cv_y = np.array(list(cv_df['Class']))

```

0.2 Apply Logistic Regression

```

In [97]: alpha = [10 ** x for x in range(-6, 3)]
print('alpha ', alpha)
cv_log_error_array = []
for i in alpha:
    print("for alpha =", i)
    clf = SGDClassifier(class_weight='balanced', alpha=i, penalty='l2', loss='log', ra
    clf.fit(train_x_onehotCoding, train_y)
    sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
    sig_clf.fit(train_x_onehotCoding, train_y)
    sig_clf_probs = sig_clf.predict_proba(cv_x_onehotCoding)
    cv_log_error_array.append(log_loss(cv_y, sig_clf_probs, labels=clf.classes_, eps=
    # to avoid rounding error while multiplying probabilities we use log-probability e
    print("Log Loss :", log_loss(cv_y, sig_clf_probs))

fig, ax = plt.subplots()
ax.plot(alpha, cv_log_error_array, c='g')
for i, txt in enumerate(np.round(cv_log_error_array, 3)):
    ax.annotate((alpha[i], str(txt)), (alpha[i], cv_log_error_array[i]))
plt.grid()
plt.title("Cross Validation Error for each alpha")
plt.xlabel("Alpha i's")
plt.ylabel("Error measure")
plt.show()

best_alpha = np.argmin(cv_log_error_array)
clf = SGDClassifier(class_weight='balanced', alpha=alpha[best_alpha], penalty='l2', l
clf.fit(train_x_onehotCoding, train_y)
sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
sig_clf.fit(train_x_onehotCoding, train_y)

```

```

predict_y = sig_clf.predict_proba(train_x_onehotCoding)
print('For values of best alpha = ', alpha[best_alpha], "The train log loss is:",log_loss)
predict_y = sig_clf.predict_proba(cv_x_onehotCoding)
print('For values of best alpha = ', alpha[best_alpha], "The cross validation log loss is:",log_loss)
predict_y = sig_clf.predict_proba(test_x_onehotCoding)
print('For values of best alpha = ', alpha[best_alpha], "The test log loss is:",log_loss)

```

```
alpha [1e-06, 1e-05, 0.0001, 0.001, 0.01, 0.1, 1, 10, 100]
```

```
for alpha = 1e-06
```

```
Log Loss : 1.5444196654896039
```

```
for alpha = 1e-05
```

```
Log Loss : 1.5264655718196316
```

```
for alpha = 0.0001
```

```
Log Loss : 1.5353204771419813
```

```
for alpha = 0.001
```

```
Log Loss : 1.453694031210995
```

```
for alpha = 0.01
```

```
Log Loss : 1.158894546919099
```

```
for alpha = 0.1
```

```
Log Loss : 1.1518830734110763
```

```
for alpha = 1
```


```
Log Loss : 1.2495390315995822
```

```
for alpha = 10
```

```
Log Loss : 1.3275090393110414
```

```
for alpha = 100
```

```
Log Loss : 1.3465194576247437
```



For values of best alpha = 0.1 The train log loss is: 0.856020140572819
 For values of best alpha = 0.1 The cross validation log loss is: 1.1518830734110763
 For values of best alpha = 0.1 The test log loss is: 1.1833367833247161

```
In [98]: clf = SGDClassifier(class_weight='balanced', alpha=alpha[best_alpha], penalty='l2', l1=0.001)
         predict_and_plot_confusion_matrix(train_x_onehotCoding, train_y, cv_x_onehotCoding, cv_y)
```


Log loss : 1.1518830734110763

Number of mis-classified points : 0.39473684210526316


----- Confusion matrix -----



----- Precision matrix (Column Sum=1) -----



----- Recall matrix (Row sum=1) -----



1 Feature Engineering

1.1 Extra Features

Trying 'number of word' and 'lenth of text' features with 'CountVectorizer'.

```
In [99]: temp_train_df = train_df.copy()
temp_train_df['text_words'] = temp_train_df['TEXT'].apply(lambda row: len(row.split(" ")))
temp_train_df['text_length'] = temp_train_df['TEXT'].str.len()

temp_cv_df = cv_df.copy()
temp_cv_df['text_words'] = temp_cv_df['TEXT'].apply(lambda row: len(row.split(" ")))
temp_cv_df['text_length'] = temp_cv_df['TEXT'].str.len()

temp_test_df = test_df.copy()
temp_test_df['text_words'] = temp_test_df['TEXT'].apply(lambda row: len(row.split(" ")))
temp_test_df['text_length'] = temp_test_df['TEXT'].str.len()

In [100]: temp_train_df = temp_train_df[temp_train_df['text_words'].notnull() ]
temp_cv_df = temp_cv_df[temp_cv_df['text_words'].notnull()]
temp_test_df = temp_test_df[temp_test_df['text_words'].notnull()]

temp_train_df = temp_train_df[temp_train_df['text_length'].notnull() ]
temp_cv_df = temp_cv_df[temp_cv_df['text_length'].notnull()]
temp_test_df = temp_test_df[temp_test_df['text_length'].notnull()]
```

```

In [101]: from sklearn_pandas import DataFrameMapper

mapper = DataFrameMapper([
    ('TEXT', CountVectorizer(ngram_range=(1,1), min_df=10)),
    ('text_words', None),
    ('text_length', None),

])

train_features = mapper.fit_transform(temp_train_df)
cv_features = mapper.transform(temp_cv_df)
test_features = mapper.transform(temp_test_df)

cv_log_error_array=[]
for i in alpha:
    clf = SGDClassifier(alpha=i, penalty='l2', loss='log', random_state=42)
    clf.fit(train_features, y_train)

    sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
    sig_clf.fit(train_features, y_train)
    predict_y = sig_clf.predict_proba(cv_features)
    cv_log_error_array.append(log_loss(y_cv, predict_y, labels=clf.classes_, eps=1e-7))
    print('For values of alpha = ', i, "The log loss is:", log_loss(y_cv, predict_y, labels=clf.classes_, eps=1e-7))


fig, ax = plt.subplots()
ax.plot(alpha, cv_log_error_array, c='g')
for i, txt in enumerate(np.round(cv_log_error_array, 3)):
    ax.annotate((alpha[i], np.round(txt, 3)), (alpha[i], cv_log_error_array[i]))
plt.grid()
plt.title("Cross Validation Error for each alpha")
plt.xlabel("Alpha i's")
plt.ylabel("Error measure")
plt.show()

best_alpha = np.argmin(cv_log_error_array)
clf = SGDClassifier(alpha=alpha[best_alpha], penalty='l2', loss='log', random_state=42)
clf.fit(train_text_feature_onehotCoding, y_train)
sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
sig_clf.fit(train_text_feature_onehotCoding, y_train)

predict_y = sig_clf.predict_proba(train_text_feature_onehotCoding)
print('For values of best alpha = ', alpha[best_alpha], "The train log loss is:", log_loss(y_train, predict_y, labels=clf.classes_, eps=1e-7))
predict_y = sig_clf.predict_proba(cv_text_feature_onehotCoding)
print('For values of best alpha = ', alpha[best_alpha], "The cross validation log loss is:", log_loss(y_cv, predict_y, labels=clf.classes_, eps=1e-7))
predict_y = sig_clf.predict_proba(test_text_feature_onehotCoding)
print('For values of best alpha = ', alpha[best_alpha], "The test log loss is:", log_loss(y_test, predict_y, labels=clf.classes_, eps=1e-7))

```

For values of alpha = 1e-06 The log loss is: 1.822160579713747
 For values of alpha = 1e-05 The log loss is: 1.8221605797145237
 For values of alpha = 0.0001 The log loss is: 1.8221605797177962
 For values of alpha = 0.001 The log loss is: 1.8221605797036327
 For values of alpha = 0.01 The log loss is: 1.8221605797138882
 For values of alpha = 0.1 The log loss is: 1.8208093577764521
 For values of alpha = 1 The log loss is: 1.8191776924723901
 For values of alpha = 10 The log loss is: 1.816933666180908
 For values of alpha = 100 The log loss is: 1.8166951485444496



For values of best alpha = 100 The train log loss is: 1.2361499752558733
 For values of best alpha = 100 The cross validation log loss is: 1.3699260961629411
 For values of best alpha = 100 The test log loss is: 1.3815368312184526

Results are not quite good. Log loss is around 1.37 which bit higher.

Let's try a tfidf vectorizer with 'no of words' and 'lenth of text' feature

```
In [108]: mapper = DataFrameMapper([
    ('TEXT', TfidfVectorizer(ngram_range=(1,2), min_df=10)),
    ('text_words', None),
    ('text_length', None),
])
```

```

train_features = mapper.fit_transform(temp_train_df)
cv_features = mapper.transform(temp_cv_df)
test_features = mapper.transform(temp_test_df)

```

```

In [109]: cv_log_error_array=[]
for i in alpha:
    clf = SGDClassifier(alpha=i, penalty='l2', loss='log', random_state=42)
    clf.fit(train_features, y_train)

    sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
    sig_clf.fit(train_features, y_train)
    predict_y = sig_clf.predict_proba(cv_features)
    cv_log_error_array.append(log_loss(y_cv, predict_y, labels=clf.classes_, eps=1e-7))
    print('For values of alpha = ', i, "The log loss is:", log_loss(y_cv, predict_y, labels=clf.classes_, eps=1e-7))

fig, ax = plt.subplots()
ax.plot(alpha, cv_log_error_array, c='g')
for i, txt in enumerate(np.round(cv_log_error_array, 3)):
    ax.annotate((alpha[i], np.round(txt, 3)), (alpha[i], cv_log_error_array[i]))
plt.grid()
plt.title("Cross Validation Error for each alpha")
plt.xlabel("Alpha i's")
plt.ylabel("Error measure")
plt.show()

best_alpha = np.argmin(cv_log_error_array)
clf = SGDClassifier(alpha=alpha[best_alpha], penalty='l2', loss='log', random_state=42)
clf.fit(train_text_feature_onehotCoding, y_train)
sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
sig_clf.fit(train_text_feature_onehotCoding, y_train)

predict_y = sig_clf.predict_proba(train_text_feature_onehotCoding)
print('For values of best alpha = ', alpha[best_alpha], "The train log loss is:", log_loss(y_train, predict_y, labels=clf.classes_, eps=1e-7))
predict_y = sig_clf.predict_proba(cv_text_feature_onehotCoding)
print('For values of best alpha = ', alpha[best_alpha], "The cross validation log loss is:", log_loss(y_cv, predict_y, labels=clf.classes_, eps=1e-7))
predict_y = sig_clf.predict_proba(test_text_feature_onehotCoding)
print('For values of best alpha = ', alpha[best_alpha], "The test log loss is:", log_loss(y_test, predict_y, labels=clf.classes_, eps=1e-7))

```


```

For values of alpha = 1e-06 The log loss is: 1.8220384848405864
For values of alpha = 1e-05 The log loss is: 1.822038484839043
For values of alpha = 0.0001 The log loss is: 1.822038484867094
For values of alpha = 0.001 The log loss is: 1.8220384848376174
For values of alpha = 0.01 The log loss is: 1.8220384848515465
For values of alpha = 0.1 The log loss is: 1.8194418728366117
For values of alpha = 1 The log loss is: 1.8179773184963033

```

For values of alpha = 10 The log loss is: 1.817733943204382

For values of alpha = 100 The log loss is: 1.816565409418684



For values of best alpha = 100 The train log loss is: 1.2361499752558733

For values of best alpha = 100 The cross validation log loss is: 1.3699260961629411

For values of best alpha = 100 The test log loss is: 1.3815368312184526

Here we can observe that 'no of words' and 'lenth of Text' are not good features. Count Vectorizer and TFIDF vectorizer both gives almost equal log loss so we can ignore those features.

A new idea is we can try "**FourGrams**". But Using fourgram require high computational power, cause features will increase tremendous. So We will use **SelectKBest** to choose 1000 important feature along with TFIDF vectorizer.

```
In [91]: # building a CountVectorizer with all the words that occurred minimum 3 times in train
text_vectorizer = TfidfVectorizer(ngram_range=(1,4),min_df=10, max_features=1000)
train_text_feature_onehotCoding = text_vectorizer.fit_transform(train_df['TEXT'])
train_text_feature_onehotCoding = SelectKBest(chi2, k=1000).fit_transform(train_text_

# getting all the feature names (words)
train_text_features= text_vectorizer.get_feature_names()

# train_text_feature_onehotCoding.sum(axis=0).A1 will sum every row and returns (1*nu
train_text_fea_counts = train_text_feature_onehotCoding.sum(axis=0).A1
```

```
# zip(list(text_features),text_fea_counts) will zip a word with its number of times i
text_fea_dict = dict(zip(list(train_text_features),train_text_fea_counts))
```

```
# taking TOP 1000 features
```

```
print("Total number of unique words in train data :", len(train_text_features))
print('train_text_feature_onehotCoding shape ==>',train_text_feature_onehotCoding.shap
```

```
Total number of unique words in train data : 1000
train_text_feature_onehotCoding shape ==> (2124, 1000)
```

```
In [92]: train_text_features[:10], text_fea_dict.get('000')
```

```
Out[92]: (['000', '05', '10', '100', '11', '12', '13', '14', '15', '16'],
14.129722873192303)
```

```
In [93]: dict_list = []
# dict_list=[] contains 9 dictoinaries each corresponds to a class
for i in range(1,10):
    cls_text = train_df[train_df['Class']==i]
    # build a word dict based on the words in that class
    dict_list.append(extract_dictionary_paddle(cls_text))
    # append it to dict_list
```

```
# dict_list[i] is build on i'th class text data
# total_dict is buid on whole training text data
total_dict = extract_dictionary_paddle(train_df)
```

```
confuse_array = []
for i in train_text_features:
    ratios = []
    max_val = -1
    for j in range(0,9):
        ratios.append((dict_list[j][i]+10 )/(total_dict[i]+90))
    confuse_array.append(ratios)
confuse_array = np.array(confuse_array)
```

```
# don't forget to normalize every feature
train_text_feature_onehotCoding = normalize(train_text_feature_onehotCoding, axis=0)
```

```
# we use the same vectorizer that was trained on train data
test_text_feature_onehotCoding = text_vectorizer.transform(test_df['TEXT'])
# don't forget to normalize every feature
test_text_feature_onehotCoding = normalize(test_text_feature_onehotCoding, axis=0)
```

```
# we use the same vectorizer that was trained on train data
```

```

cv_text_feature_onehotCoding = text_vectorizer.transform(cv_df['TEXT'])
# don't forget to normalize every feature
cv_text_feature_onehotCoding = normalize(cv_text_feature_onehotCoding, axis=0)

In [94]: train_gene_var_onehotCoding = hstack((train_gene_feature_onehotCoding, train_variation_fe
test_gene_var_onehotCoding = hstack((test_gene_feature_onehotCoding, test_variation_fe
cv_gene_var_onehotCoding = hstack((cv_gene_feature_onehotCoding, cv_variation_feature_o

train_x_onehotCoding = hstack((train_gene_var_onehotCoding, train_text_feature_onehotC
train_y = np.array(list(train_df['Class'])))

test_x_onehotCoding = hstack((test_gene_var_onehotCoding, test_text_feature_onehotCod
test_y = np.array(list(test_df['Class'])))

cv_x_onehotCoding = hstack((cv_gene_var_onehotCoding, cv_text_feature_onehotCoding)).
cv_y = np.array(list(cv_df['Class'])))

In [95]: alpha = [10 ** x for x in range(-6, 3)]
print('alpha ', alpha)
cv_log_error_array = []
for i in alpha:
    print("for alpha =", i)
    clf = SGDClassifier(class_weight='balanced', alpha=i, penalty='l2', loss='log', r
    clf.fit(train_x_onehotCoding, train_y)
    sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
    sig_clf.fit(train_x_onehotCoding, train_y)
    sig_clf_probs = sig_clf.predict_proba(cv_x_onehotCoding)
    cv_log_error_array.append(log_loss(cv_y, sig_clf_probs, labels=clf.classes_, eps=
    # to avoid rounding error while multiplying probabilities we use log-probability e
    print("Log Loss :", log_loss(cv_y, sig_clf_probs))

fig, ax = plt.subplots()
ax.plot(alpha, cv_log_error_array, c='g')
for i, txt in enumerate(np.round(cv_log_error_array, 3)):
    ax.annotate((alpha[i], str(txt)), (alpha[i], cv_log_error_array[i]))
plt.grid()
plt.title("Cross Validation Error for each alpha")
plt.xlabel("Alpha i's")
plt.ylabel("Error measure")
plt.show()

best_alpha = np.argmin(cv_log_error_array)
clf = SGDClassifier(class_weight='balanced', alpha=alpha[best_alpha], penalty='l2', l
clf.fit(train_x_onehotCoding, train_y)
sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
sig_clf.fit(train_x_onehotCoding, train_y)

```



```

predict_y = sig_clf.predict_proba(train_x_onehotCoding)
print('For values of best alpha = ', alpha[best_alpha], "The train log loss is:",log_1)
predict_y = sig_clf.predict_proba(cv_x_onehotCoding)
print('For values of best alpha = ', alpha[best_alpha], "The cross validation log loss is:",log_2)
predict_y = sig_clf.predict_proba(test_x_onehotCoding)
print('For values of best alpha = ', alpha[best_alpha], "The test log loss is:",log_3)

```

```
alpha [1e-06, 1e-05, 0.0001, 0.001, 0.01, 0.1, 1, 10, 100]
```

```
for alpha = 1e-06
```

```
Log Loss : 1.212707277334664
```

```
for alpha = 1e-05
```

```
Log Loss : 1.153533111121014
```

```
for alpha = 0.0001
```

```
Log Loss : 1.0758569532173068
```

```
for alpha = 0.001
```

```
Log Loss : 1.073555694140625
```

```
for alpha = 0.01
```

```
Log Loss : 1.176356512596208
```

```
for alpha = 0.1
```

```
Log Loss : 1.7026559455757415
```

```
for alpha = 1
```


```
Log Loss : 1.8511449207433004
```

```
for alpha = 10
```

```
Log Loss : 1.8683802973850192
```

```
for alpha = 100
```

```
Log Loss : 1.8705021164147495
```



For values of best alpha = 0.001 The train log loss is: 0.7178898683114381
 For values of best alpha = 0.001 The cross validation log loss is: 1.073555694140625
 For values of best alpha = 0.001 The test log loss is: 1.025731799504609

```
In [96]: clf = SGDClassifier(class_weight='balanced', alpha=alpha[best_alpha], penalty='l2', l1=0.001)
         predict_and_plot_confusion_matrix(train_x_onehotCoding, train_y, cv_x_onehotCoding, cv_y)
```


Log loss : 1.073555694140625

Number of mis-classified points : 0.39285714285714285


----- Confusion matrix -----



----- Precision matrix (Column Sum=1) -----



----- Recall matrix (Row sum=1) -----



It can observe that log is quite low but yet it is not under 1. But it should notice that misclassified points percentage is increasing which is bad sign. RED ALERT.

We can try TDIDF and countvectorizer again it require higher computation power so we will limit ourselves to 1000 max features only.

```
In [100]: from sklearn_pandas import DataFrameMapper
```

```
mapper = DataFrameMapper([
    ('TEXT', CountVectorizer(ngram_range=(1,4), min_df=10, max_features=1000)),
    ('TEXT', TfidfVectorizer(ngram_range=(1,4), min_df=10, max_features=1000)),
])
```

```
train_features = mapper.fit_transform(train_df)
cv_features = mapper.transform(cv_df)
test_features = mapper.transform(test_df)
```

```
cv_log_error_array=[]
for i in alpha:
    clf = SGDClassifier(alpha=i, penalty='l2', loss='log', random_state=42)
    clf.fit(train_features, y_train)

    sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
    sig_clf.fit(train_features, y_train)
    predict_y = sig_clf.predict_proba(cv_features)
    cv_log_error_array.append(log_loss(y_cv, predict_y, labels=clf.classes_, eps=1e-7))
```

```

        print('For values of alpha = ', i, "The log loss is:",log_loss(y_cv, predict_y,

fig, ax = plt.subplots()
ax.plot(alpha, cv_log_error_array,c='g')
for i, txt in enumerate(np.round(cv_log_error_array,3)):
    ax.annotate((alpha[i],np.round(txt,3)), (alpha[i],cv_log_error_array[i]))
plt.grid()
plt.title("Cross Validation Error for each alpha")
plt.xlabel("Alpha i's")
plt.ylabel("Error measure")
plt.show()

best_alpha = np.argmin(cv_log_error_array)
clf = SGDClassifier(alpha=alpha[best_alpha], penalty='l2', loss='log', random_state=
clf.fit(train_text_feature_onehotCoding, y_train)
sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
sig_clf.fit(train_text_feature_onehotCoding, y_train)


predict_y = sig_clf.predict_proba(train_text_feature_onehotCoding)
print('For values of best alpha = ', alpha[best_alpha], "The train log loss is:",log
predict_y = sig_clf.predict_proba(cv_text_feature_onehotCoding)
print('For values of best alpha = ', alpha[best_alpha], "The cross validation log lo
predict_y = sig_clf.predict_proba(test_text_feature_onehotCoding)
print('For values of best alpha = ', alpha[best_alpha], "The test log loss is:",log_

```

```

For values of alpha = 1e-06 The log loss is: 1.8304997567764278
For values of alpha = 1e-05 The log loss is: 1.8304997567764278
For values of alpha = 0.0001 The log loss is: 1.3526187637425466
For values of alpha = 0.001 The log loss is: 1.323785870393301
For values of alpha = 0.01 The log loss is: 1.3366673693797064
For values of alpha = 0.1 The log loss is: 1.3289461442950286
For values of alpha = 1 The log loss is: 1.309394611086164
For values of alpha = 10 The log loss is: 1.2431744140109304
For values of alpha = 100 The log loss is: 1.3277691185564866

```



For values of best alpha = 10 The train log loss is: 1.78441584487254

For values of best alpha = 10 The cross validation log loss is: 2.1767942622361502


For values of best alpha = 10 The test log loss is: 1.9985137946444134

```
In [101]: clf = SGDClassifier(class_weight='balanced', alpha=alpha[best_alpha], penalty='l2',
    predict_and_plot_confusion_matrix(train_x_onehotCoding, train_y, cv_x_onehotCoding, cv_y)
```

Log loss : 1.8683802973850192

Number of mis-classified points : 0.7236842105263158


----- Confusion matrix -----



----- Precision matrix (Column Sum=1) -----



----- Recall matrix (Row sum=1) -----



It was the worst model among all we tried. Probably because alone countvectorizer has 53k features and we are dealing with only 1k.

2 Conclusion

```
In [70]: import pandas as pd
          from prettytable import PrettyTable
```

```
print('\t\t\t\t Model Comparision ')
print('\n')
```

```
x = PrettyTable()
x.field_names = ['Metric','RM','NB','KNN', 'LR-bal', 'LR-w/o', 'LinearSVM', 'RF-res']

x.add_row(["Tr Loss ",2.50,0.52342, 0.50799, 0.43916,0.42938, 0.47877,0.85274,0.05408]
x.add_row(["Test Loss ",2.50,1.18371, 1.00464,0.97336,1.01902,1.01222,1.17696,1.28650]
x.add_row(["Mis-class ",0.0,0.37781, 0.36278, 0.34774,0.34962,0.33646,0.40037,0.49624]

print('\n')
print(x)
```

Model Comparision

Metric	RM	NB	KNN	LR-bal	LR-w/o	LinearSVM	RF-res	RF-one hot
Tr Loss	2.5	0.52342	0.50799	0.43916	0.42938	0.47877	0.85274	0.05408
Test Loss	2.5	1.18371	1.00464	0.97336	1.01902	1.01222	1.17696	1.2865
Mis-class	0.0	0.37781	0.36278	0.34774	0.34962	0.33646	0.40037	0.49624

```
In [71]: x1 = PrettyTable()
```

```
x1.field_names = ['Metric','Vectorizer','Feature','TrLoss', 'TestLoss','Mis-class']
x1.add_row(["LR ','Countvect','Bigram' ,0.85602, 1.15188, 0.39473])
x1.add_row(["LR ','Countvect','no_of words +text len' ,1.23614, 1.36992, 0.38873])
x1.add_row(["LR ','TFIDF','no_of words +text len' ,1.23614, 1.36992, 0.38873])
x1.add_row(["LR ','TFIDF','fourGrams' ,0.71788, 1.07355, 0.3928])
x1.add_row(["LR ','TFIDF + CountVect','--' ,1.78441, 2.176795, 0.72368])

print(x1)
```

Metric	Vectorizer	Feature	TrLoss	TestLoss	Mis-class
LR	Countvect	Bigram	0.85602	1.15188	0.39473
LR	Countvect	no_of words +text len	1.23614	1.36992	0.38873
LR	TFIDF	no_of words +text len	1.23614	1.36992	0.38873

LR	TFIDF	fourGrams	0.71788	1.07355	0.3928	
LR	TFIDF + CountVect	--	1.78441	2.176795	0.72368	
+-----+-----+-----+-----+-----+-----+						

- We can say that Logistic regression model with **SelectBestK** is 1000 is far better model which test log loss under 1.000.
- We tried feature engineering, TFIDF with fourgram seems good, but due to computational power consumption we limit ourselves to some set of points we can perform the same with better computational power or more optimized way.